EE-219 Project 1
Shrey Agarwal (004943082)
Varun Saboo (505028591)
February 12, 2018

Project 2: Clustering

Objective –
To find proper representations of the data, s.t. the clustering is efficient and gives out reasonable results.
To perform K-means clustering on the dataset, and evaluate the performance of the clustering.
To try different preprocess methods which may increase the performance of the clustering.

Dataset –
We work with “20 Newsgroups” dataset. It is a collection of approximately 20,000 documents, partitioned (nearly) evenly across 20 different newsgroups, each corresponding to a different topic. Each topic can be viewed as a “class”.

TASK 1:

We perform TF-IDF on the documents by performing certain preprocessing steps such as stemming, word tokenizing, special characters removal and stopwords removal.
With min_df = 3, the dimension of the TF-IDF matrix is (7882, 16564), where 7882 are the total number of documents in 8 classes and 16564 are the unique tokens identified.

TASK 2:

After applying k-means clustering on the above dataset, we got the following results.
With number of clusters = 2:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	109.23s	7484	0.426	0.460	0.443	0.444	0.426

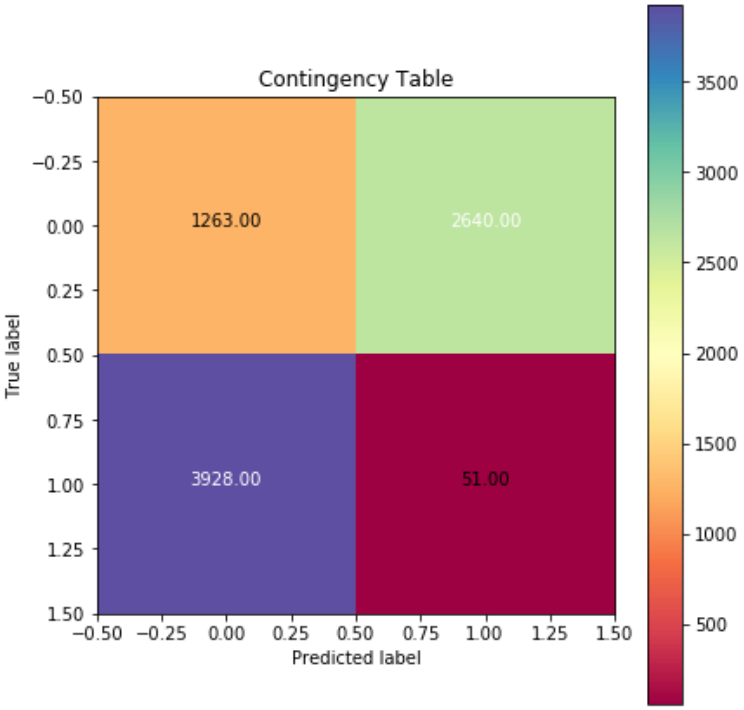

TASK 3:

Preprocess the data:

Dimensionality reduction:

a. LSI (Latent Semantic Indexing):
   After performing LSI on the TF-IDF matrix, following are the obtained results:

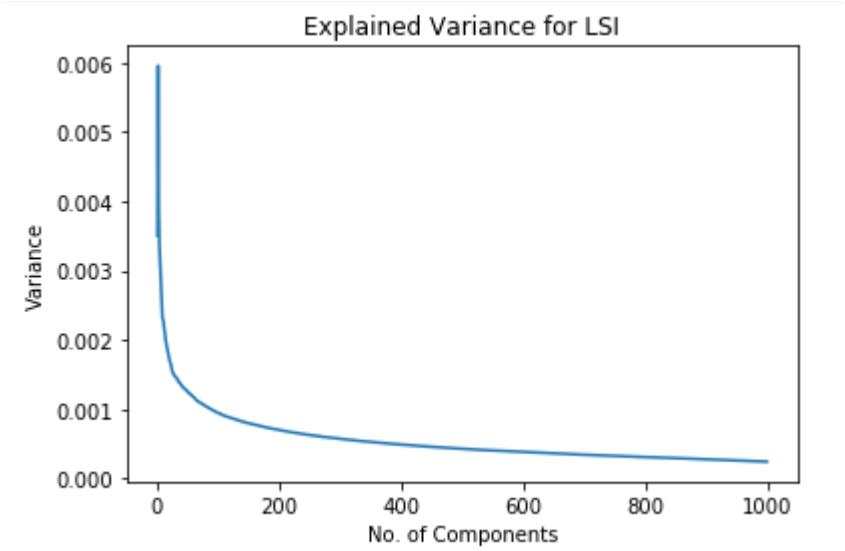

Figure 1: We can observe that the variance decreases as the no. of components increase.

Model with 1 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	0.24s	9	0.018	0.018	0.018	0.025	0.018

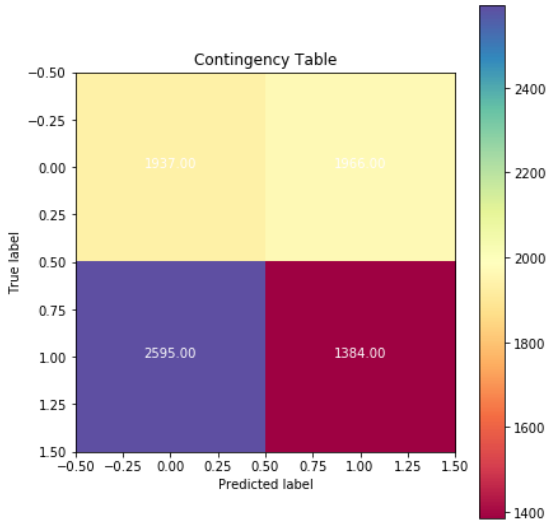

Model with 2 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	0.40s	41	0.419	0.446	0.432	0.454	0.419

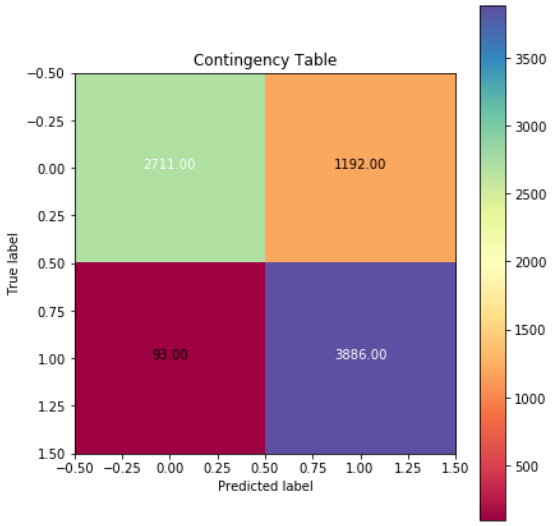

Model with 3 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	0.23s	71	0.412	0.440	0.425	0.445	0.412

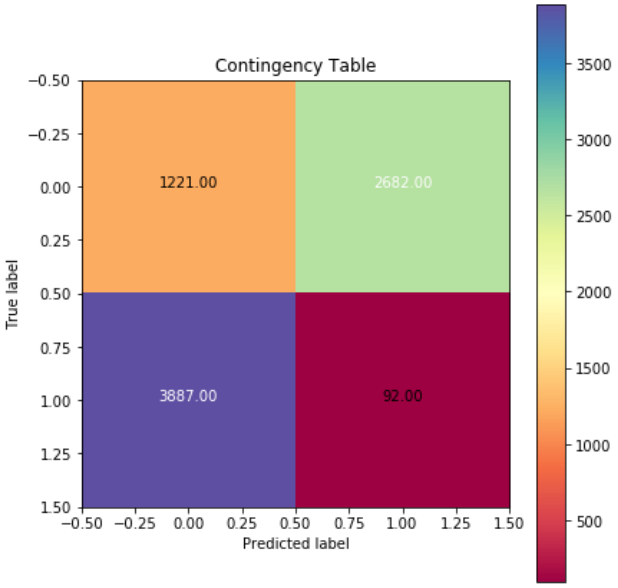

Model with 5 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	0.24s	119	0.407	0.442	0.424	0.424	0.407

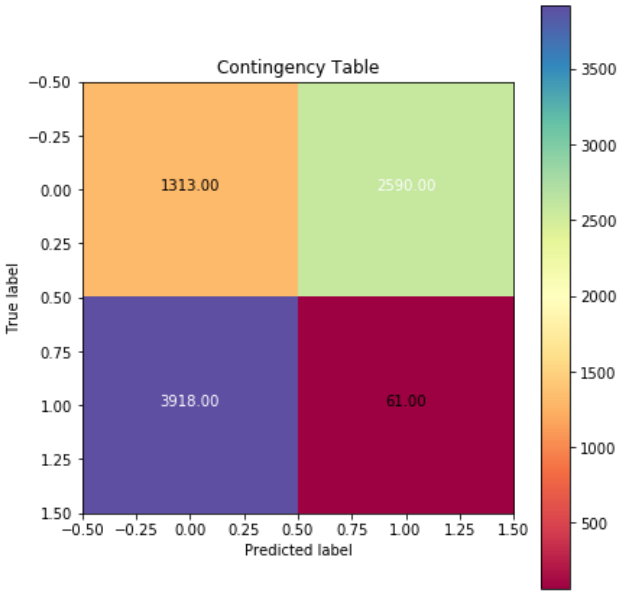

Model with 10 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	0.56s	216	0.430	0.460	0.444	0.458	0.430

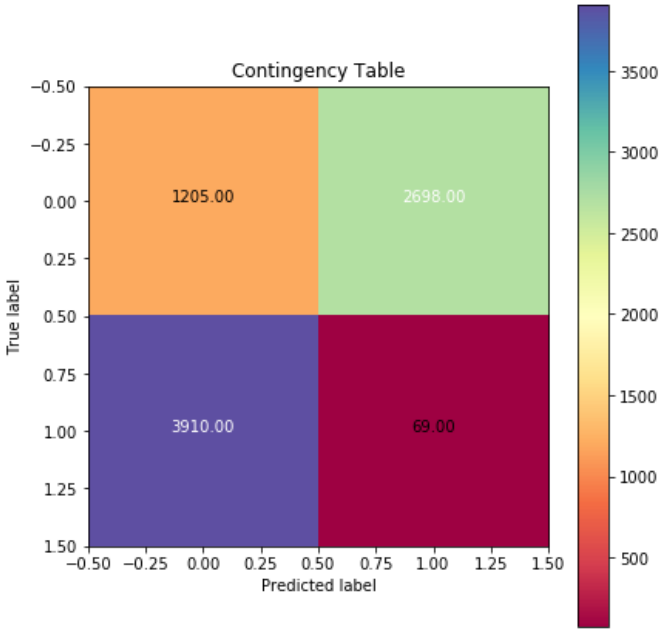

Model with 20 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	0.74s	363	0.423	0.455	0.439	0.445	0.423

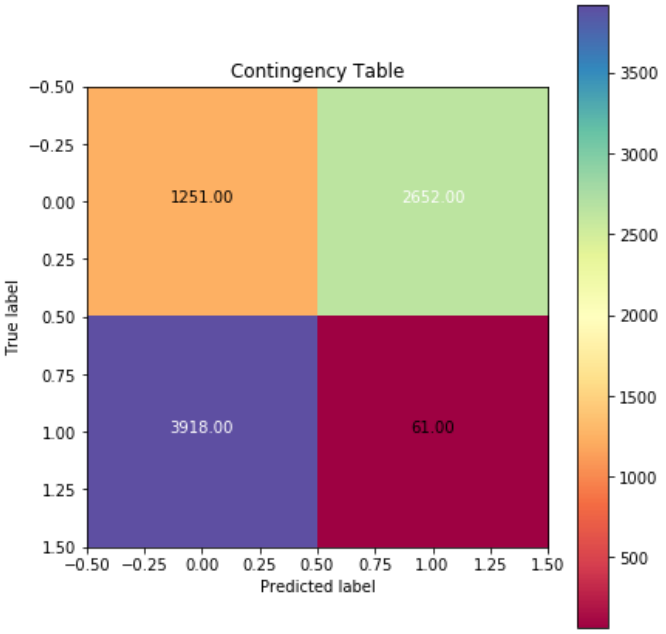

Model with 50 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	1.34s	684	0.420	0.455	0.437	0.436	0.420

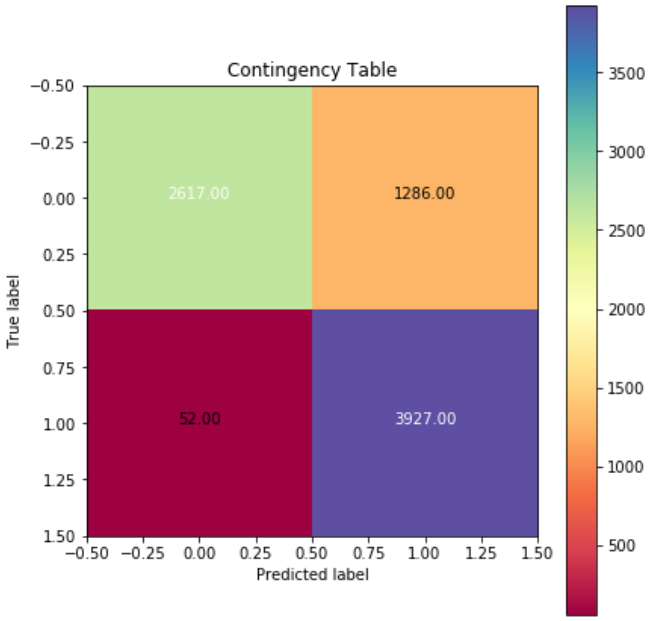

Model with 100 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	1.64s	1090	0.425	0.460	0.442	0.443	0.425

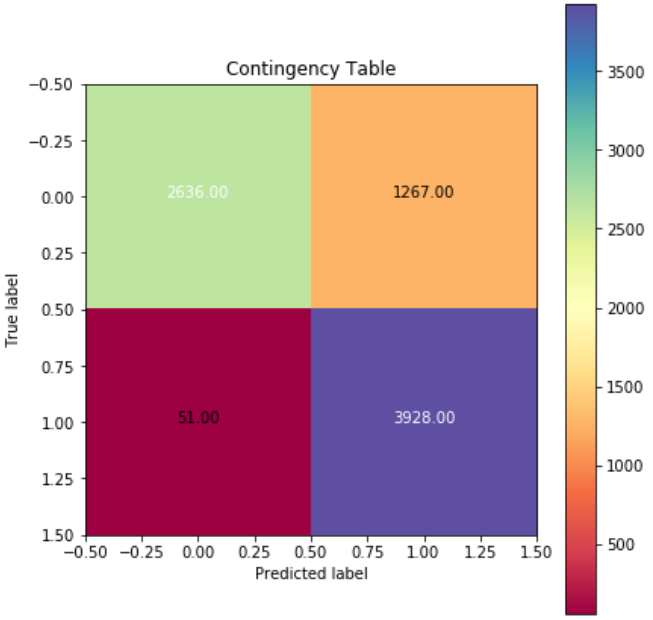

Model with 300 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	5.07s	2172	0.418	0.454	0.436	0.431	0.418

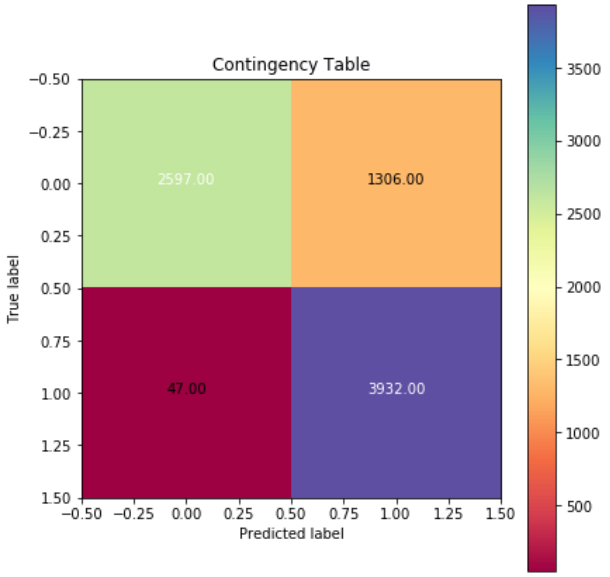

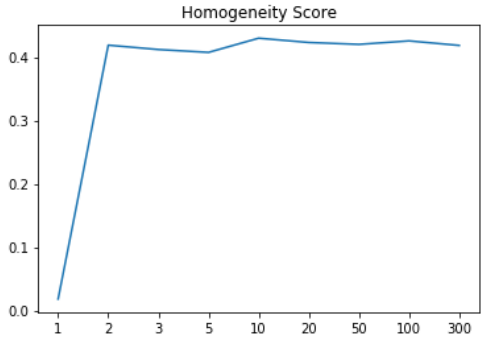

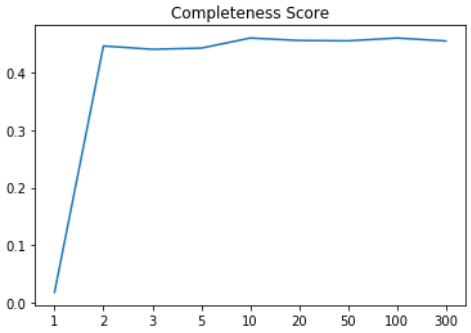

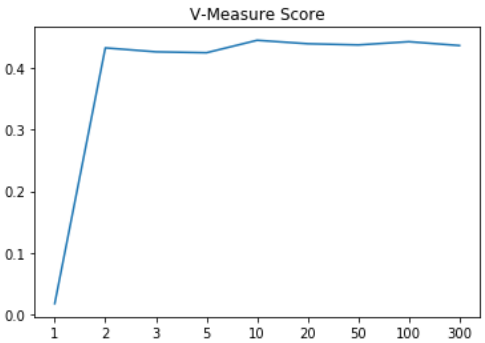

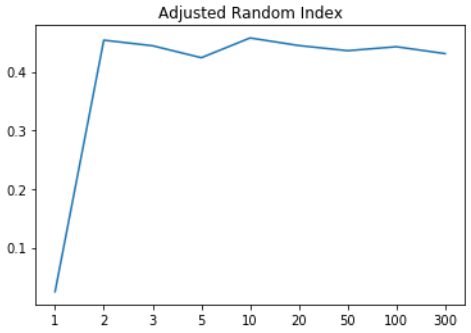

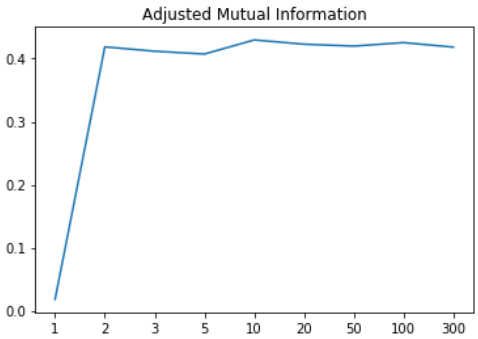

As we can observe from the contingency matrix as well as the 4 measures above, the best r value for LSI is 10. i.e, 10 components are best suitable to represent each document and cluster with maximum accuracy.

b. NMF (Non-Negative Matrix Factorization):
After performing NMF on the TF-IDF matrix, following are the obtained results:

Model with 1 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	0.14s	0	0.018	0.018	0.018	0.025	0.018

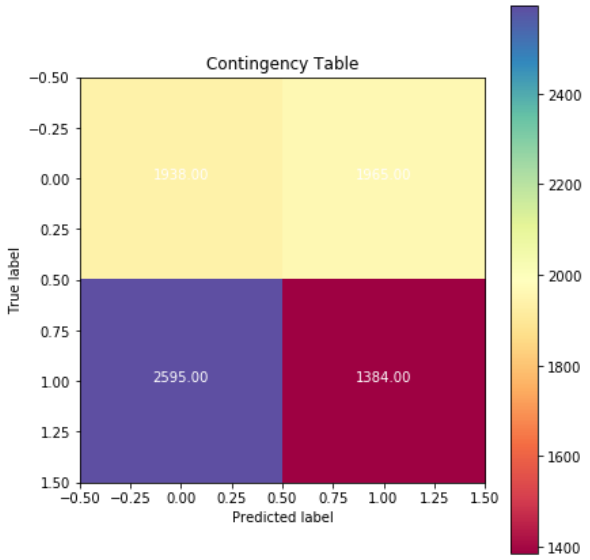

Model with 2 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	0.22s	1	0.409	0.440	0.424	0.434	0.409

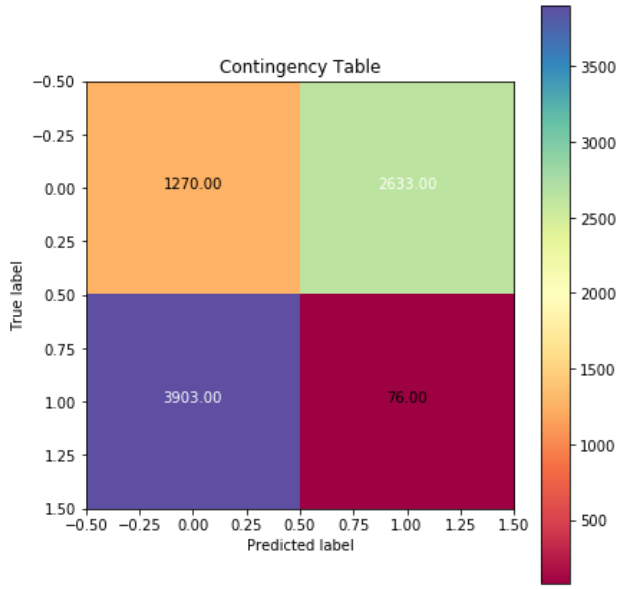

Model with 3 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	0.15s	3	0.384	0.417	0.400	0.406	0.384

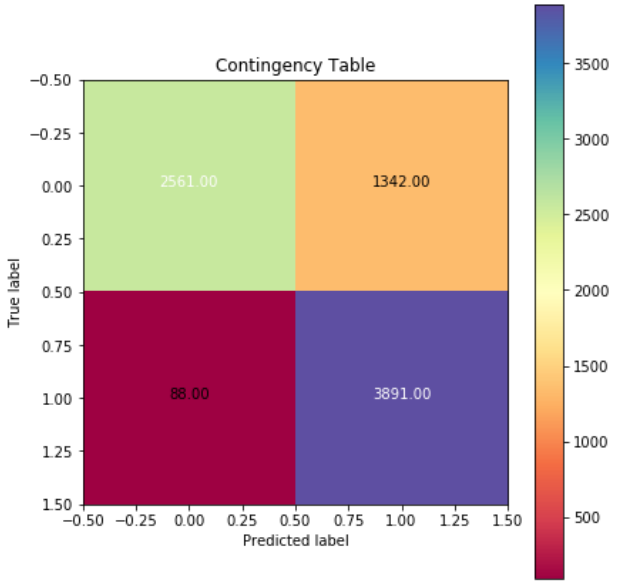

Model with 5 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	0.17s	8	0.365	0.388	0.376	0.410	0.365

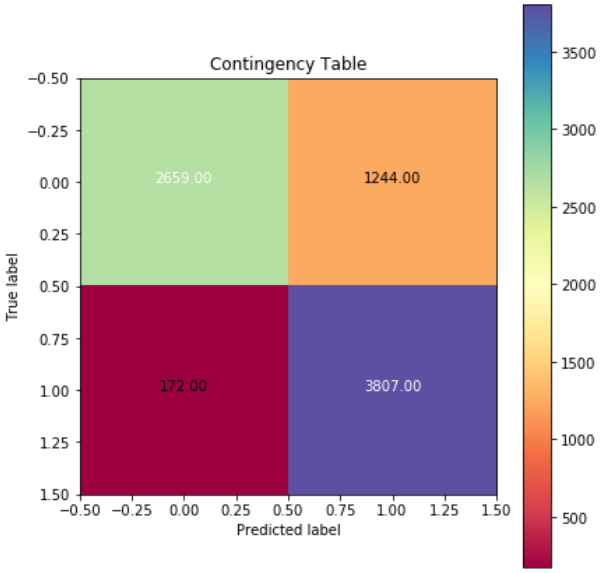

Model with 10 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	0.48s	18	0.046	0.058	0.051	0.043	0.046

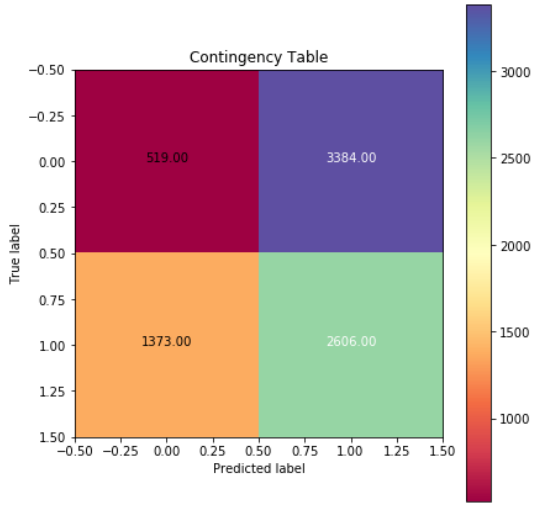

Model with 20 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	0.70s	35	0.048	0.165	0.074	0.008	0.048

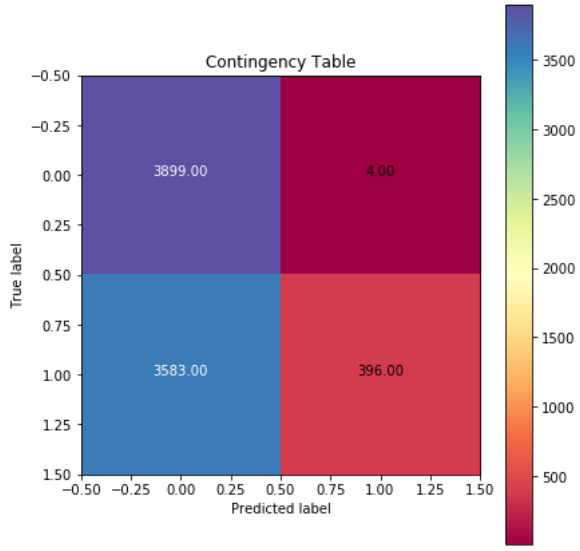

Model with 50 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	0.81s	81	0.008	0.027	0.012	0.003	0.008

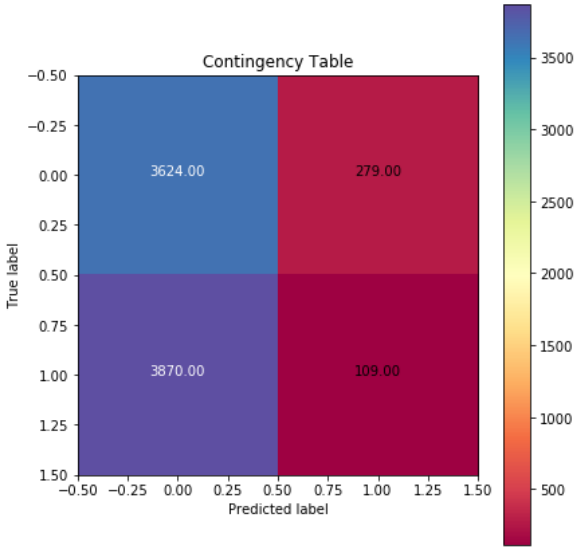

Model with 100 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	1.55s	159	0.000	0.000	0.000	-0.000	-0.000

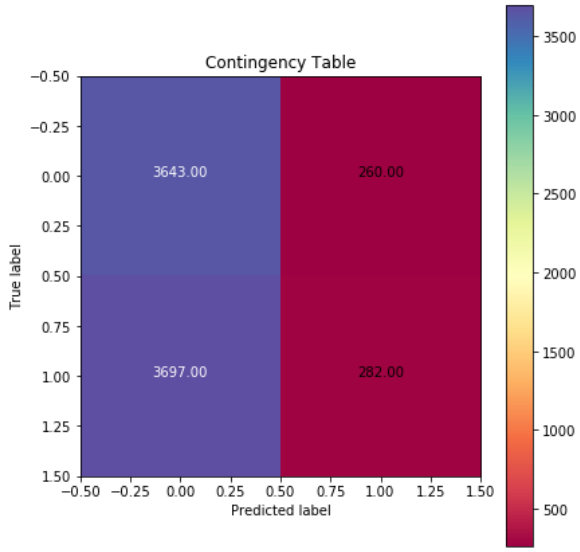

```
Model with 300 components:
_____________________________________________________________________________
init            time    inertia homo    comp    v-meas  ARI     AMI
_____________________________________________________________________________
k-means++       2.92s   420     0.010   0.048   0.016   0.002   0.009
```

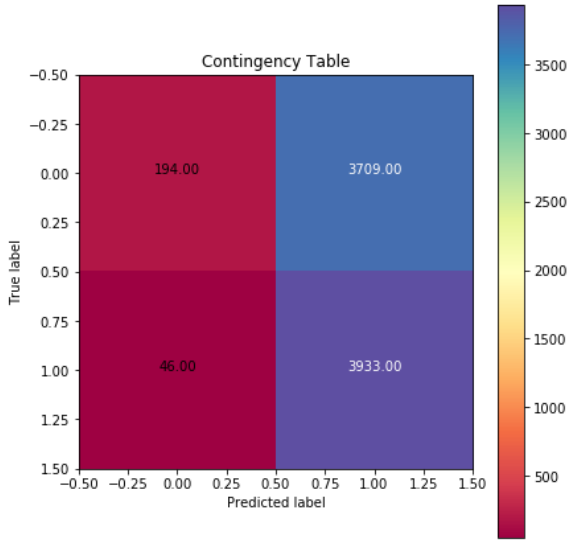

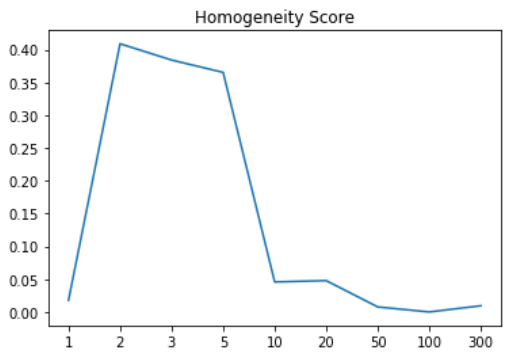

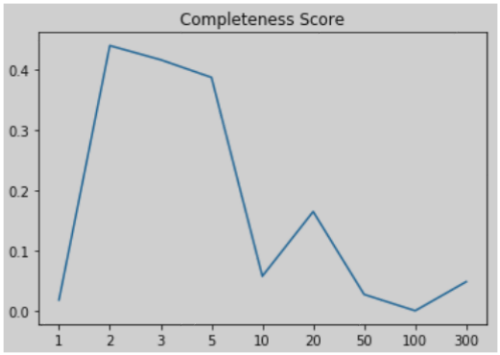

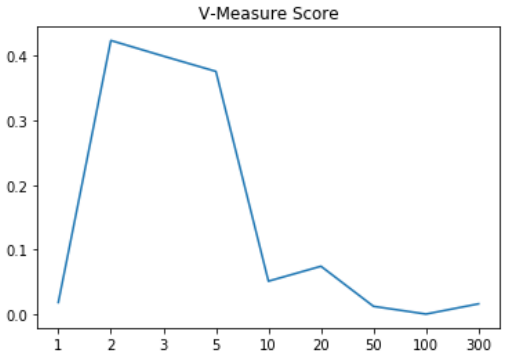

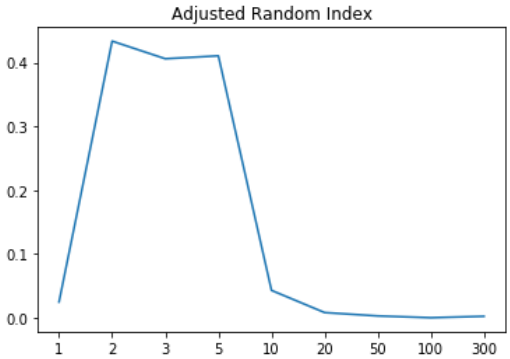

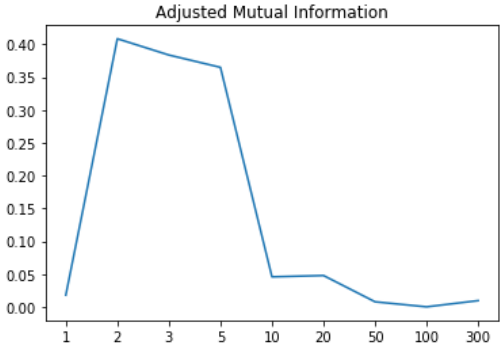

As we can observe, the best r value is at r = 2. i.e, only 2 components/features are more than enough to best cluster the documents in two classes with maximum

accuracy.

The non-monotonic behavior can be explained by the following reasoning.
The Curse of Dimensionality:
The distance to the nearest neighbor and the distance to the farthest neighbor tend to converge as the dimension increases. On the other hand, increasing the dimension of data representation provides intricate details which helps in distinguishing each data point. Hence, a proper balance needs to be found to represent data in best dimension space. Here, we observed that the clustering improved performance as we increase dimension but started performing worse for higher dimensions. Hence, it exhibits the non-monotonic behavior.

TASK 4:

Visualize the performance:

a. We found that,
   r=5 works best in case of LSI, and
   r=2 works best in case of NMF.
   Hence, following are the clusters with their decision boundary for the two methods.

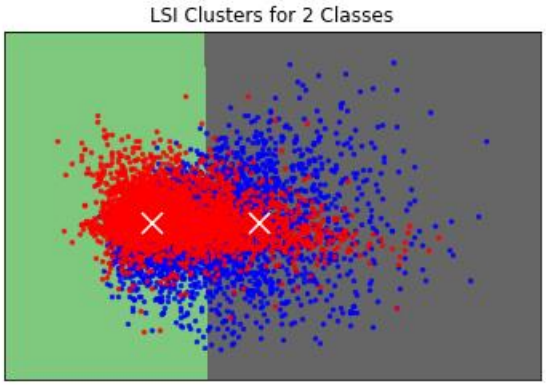

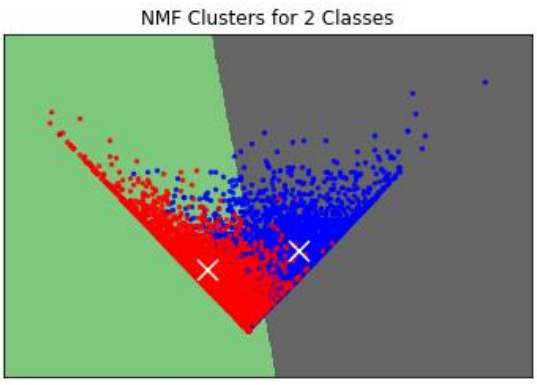

b. 3 Methods:
   i. Normalizing:
   We performed scaling using sklearn “scale” to get unit variance and then normalized the data. Following are the results for LSI and NMF:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	0.56s	9	0.018	0.021	0.019	0.021	0.018

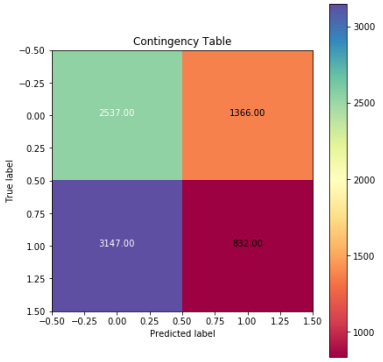

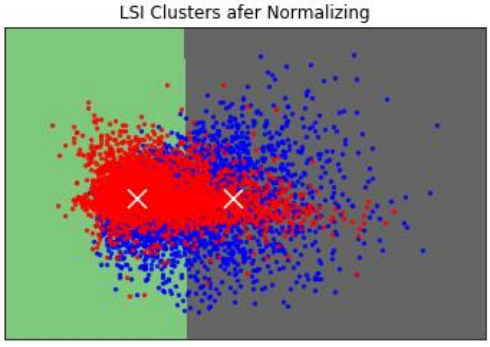

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	0.17s	1	0.450	0.468	0.459	0.510	0.450

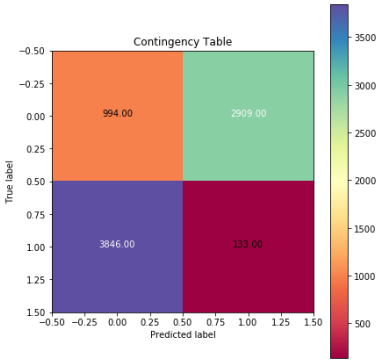

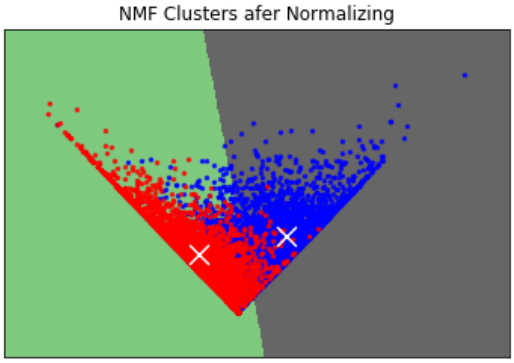

As we can observe, the performance has improved from 83.69% to 85.7% which is a good improvement.

ii. Logarithmic Transformation:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	0.19s	12022	0.467	0.469	0.468	0.567	0.467

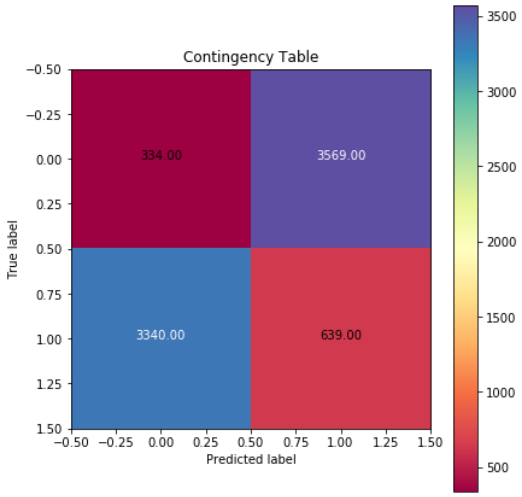

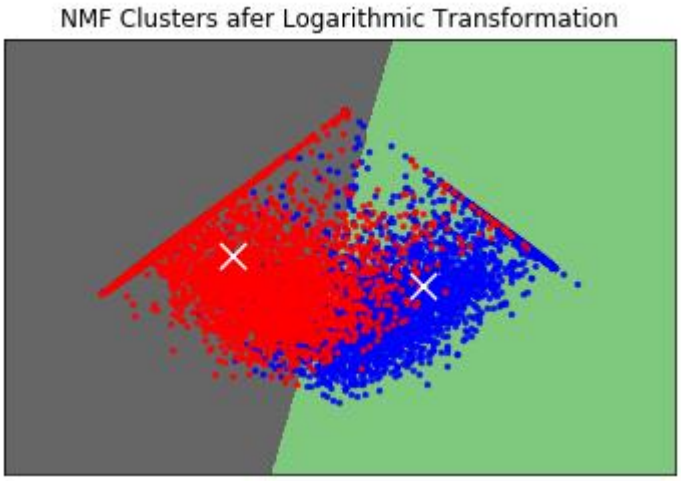

As we can observe, the performance has improved from 83.69% to 87.65% which is a very good improvement.

iii. Combination

1. Normalization and then Log Transformation

We perform normalization and then apply log transformation on NMF data to get the following results:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	0.15s	9081	0.495	0.495	0.495	0.602	0.495

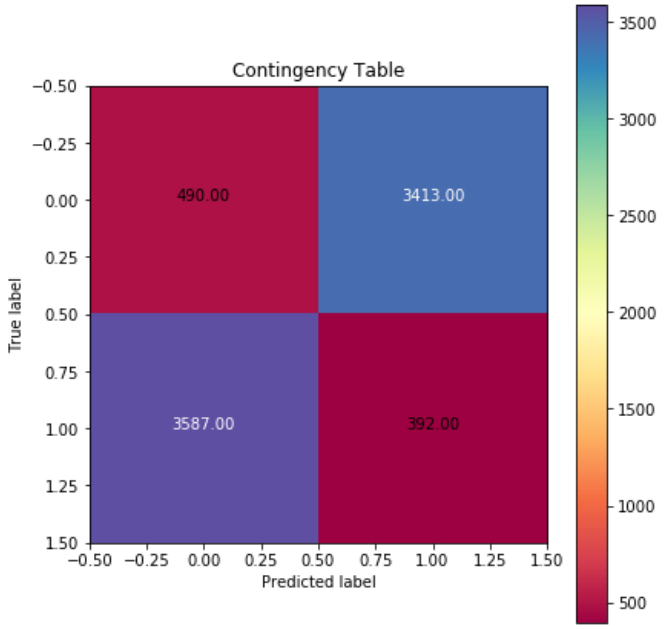

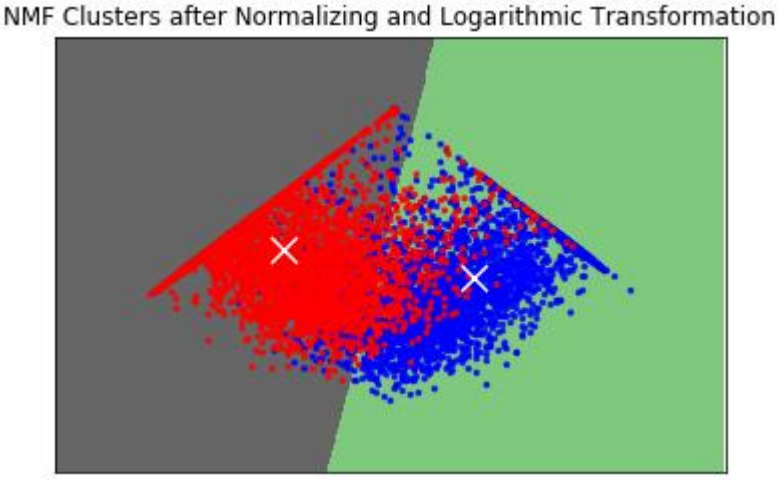

As we can observe, the performance has improved from 83.69% to 88.80% which is higher than normalization or log transformation separately. Infact, it is the highest and best combination so far.

2. Log Transformation and then Normalization

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	0.18s	0	0.489	0.490	0.489	0.597	0.489

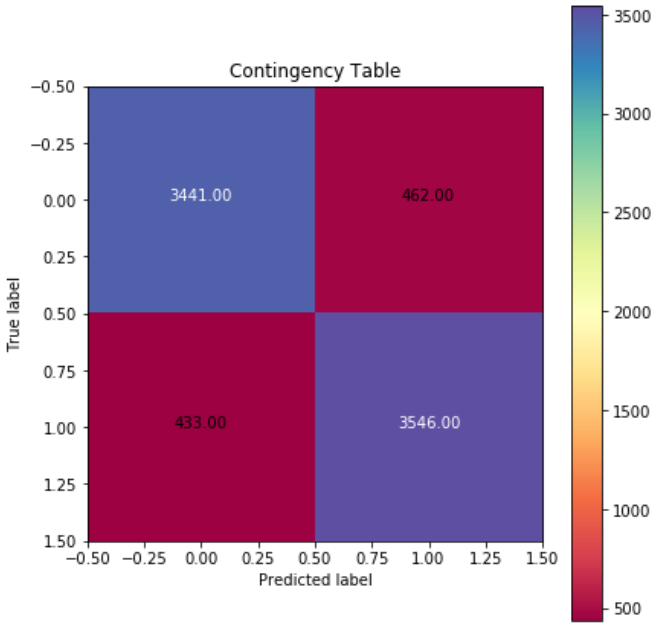

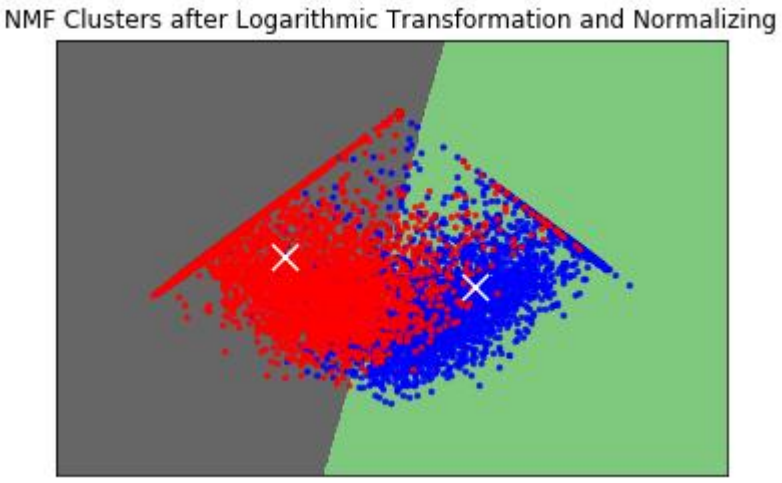

As we can observe, the performance has improved from 83.69% to 88.64% which is higher than normalization or log transformation separately.

TASK 5: Expand to 20 categories

We follow the same workflow as above with all 20 categories and produce following results and observations:

1. We perform TF-IDF on the documents by performing certain preprocessing steps such as stemming, word tokenizing, special characters removal and stopwords removal.
   With min_df = 3, the dimension of the TF-IDF matrix is (18846, 33158), where 18846 are the total number of documents in 20 categories and 33158 are the unique tokens identified.
2. After applying k-means clustering on the above dataset, we got the following results.
   With number of clusters = 2:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	570.00s	17686	0.246	0.306	0.273	0.050	0.243

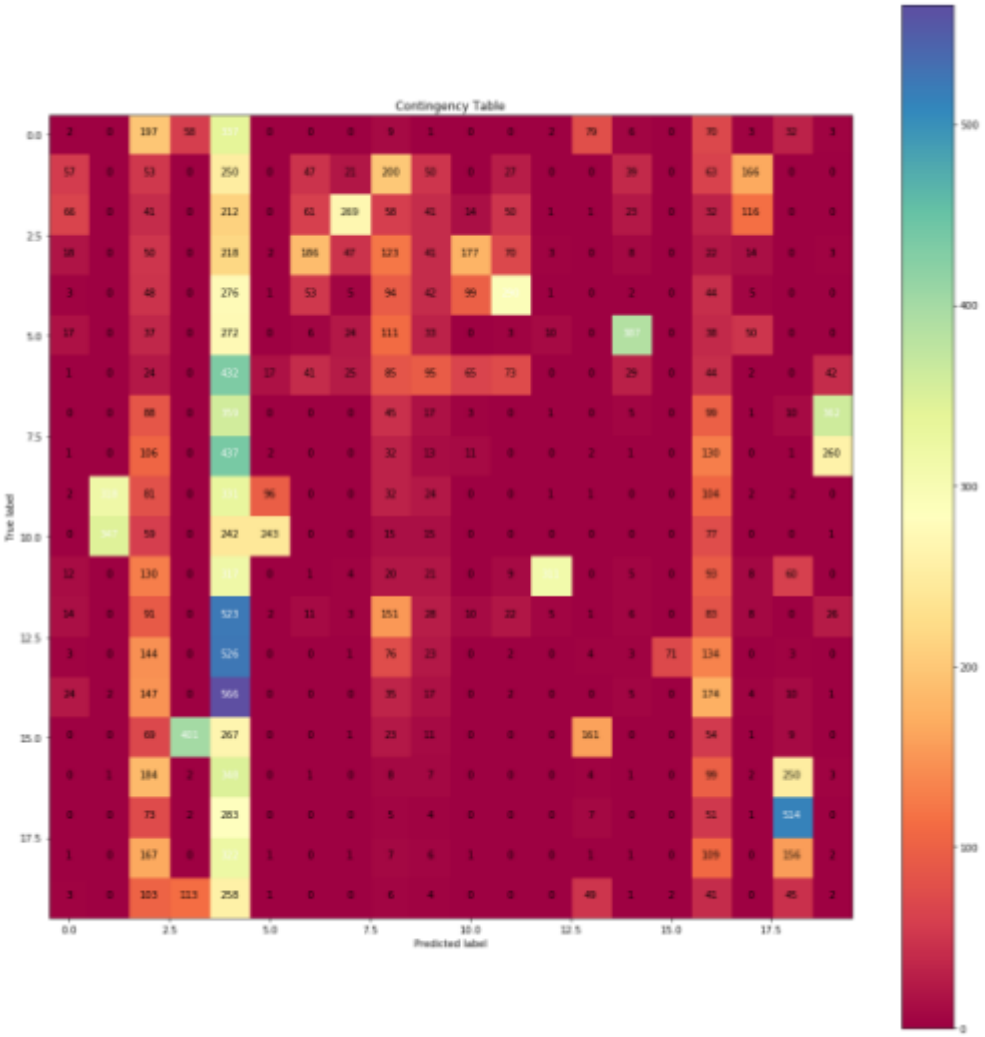

3. Preprocess the data:
   Dimensionality reduction:
   a. LSI (Latent Semantic Indexing):

   After performing LSI on the TF-IDF matrix, following are the obtained results:

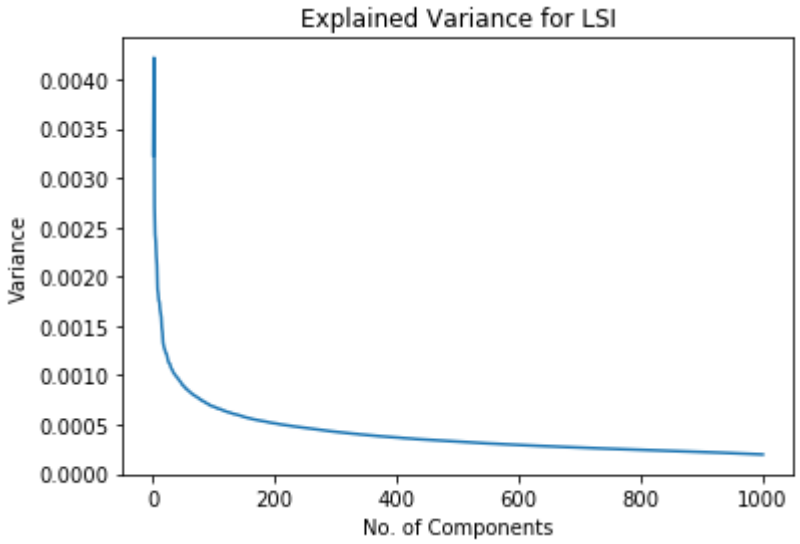

Figure 2: We can observe that the variance decreases as the no. of components increase. Interestingly, the variance is even more lesser than previous 8 categories.

Model with 1 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	6.53s	0	0.013	0.015	0.014	0.003	0.010

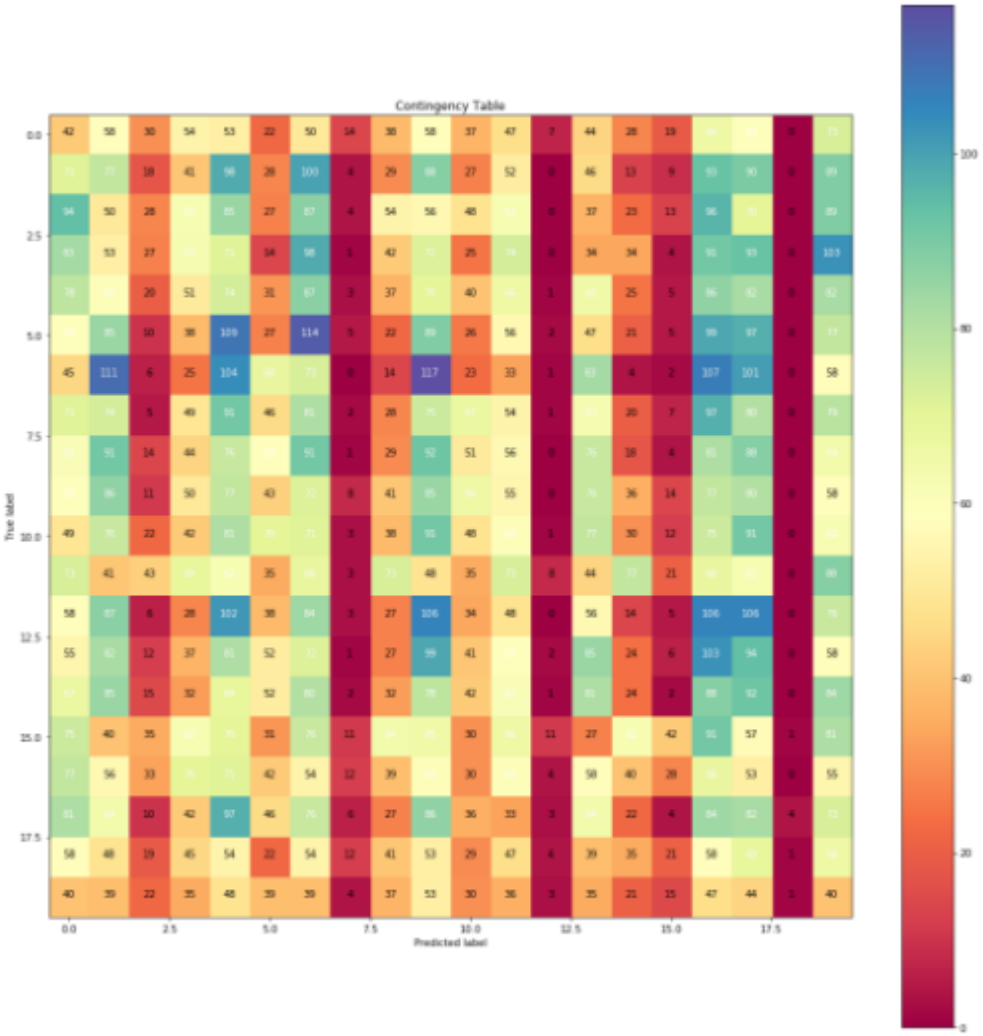

Model with 2 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	10.54s	9	0.152	0.162	0.157	0.043	0.150

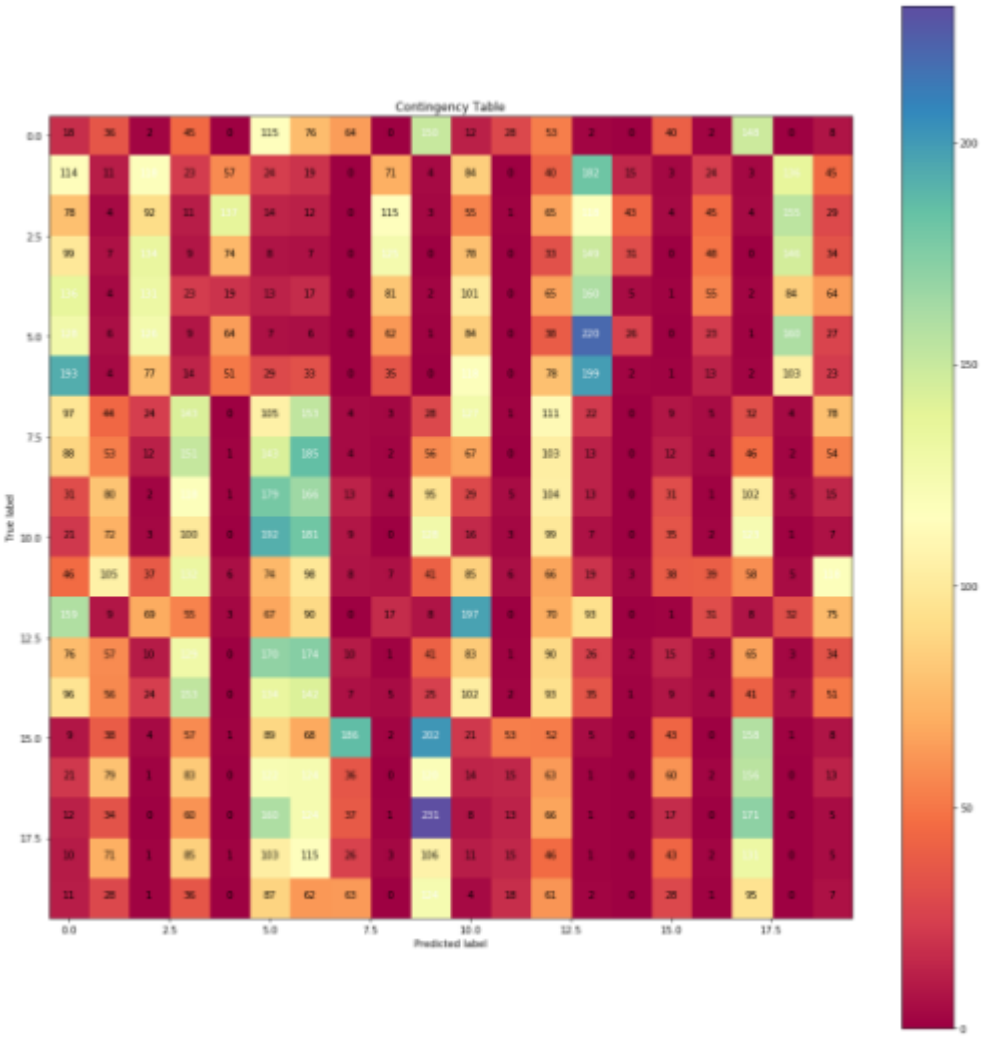

Model with 3 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	10.64s	23	0.201	0.217	0.209	0.064	0.199

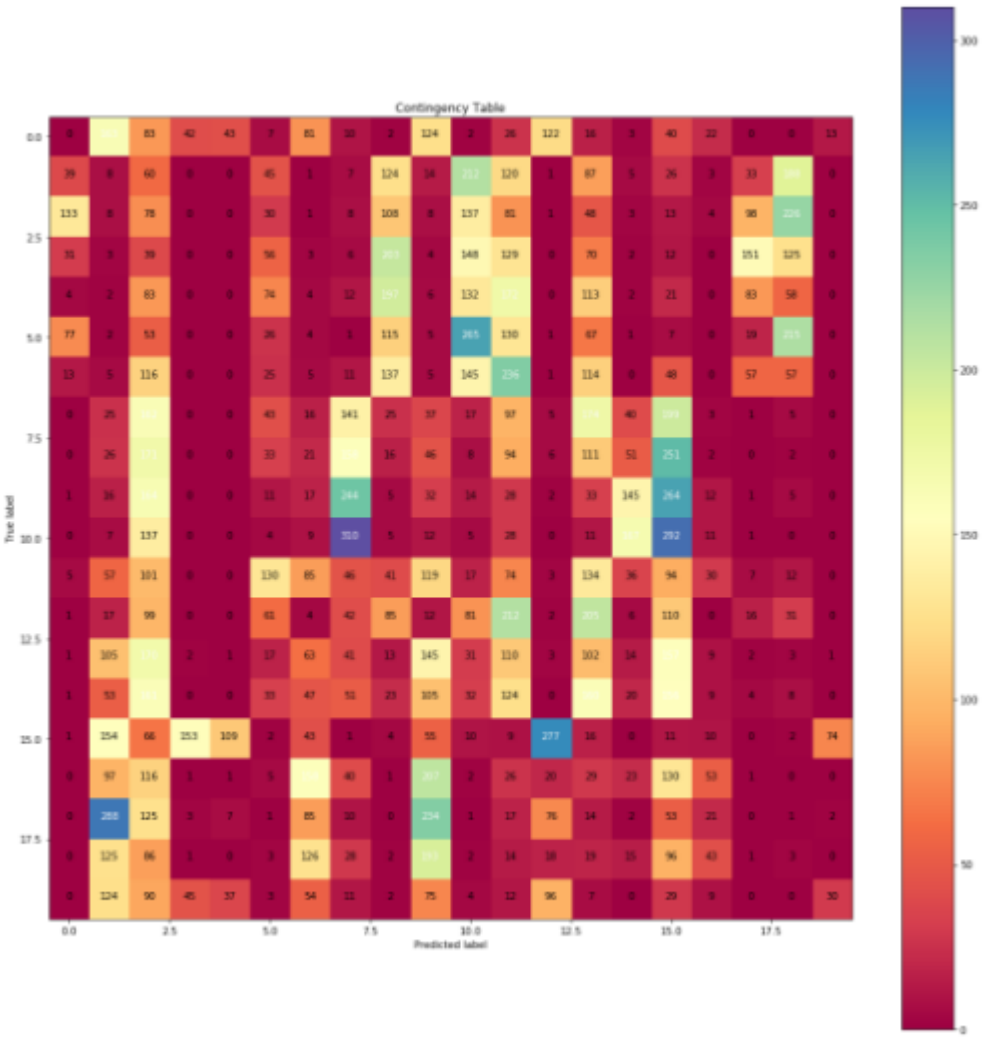

Model with 5 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	12.64s	61	0.199	0.221	0.209	0.061	0.197

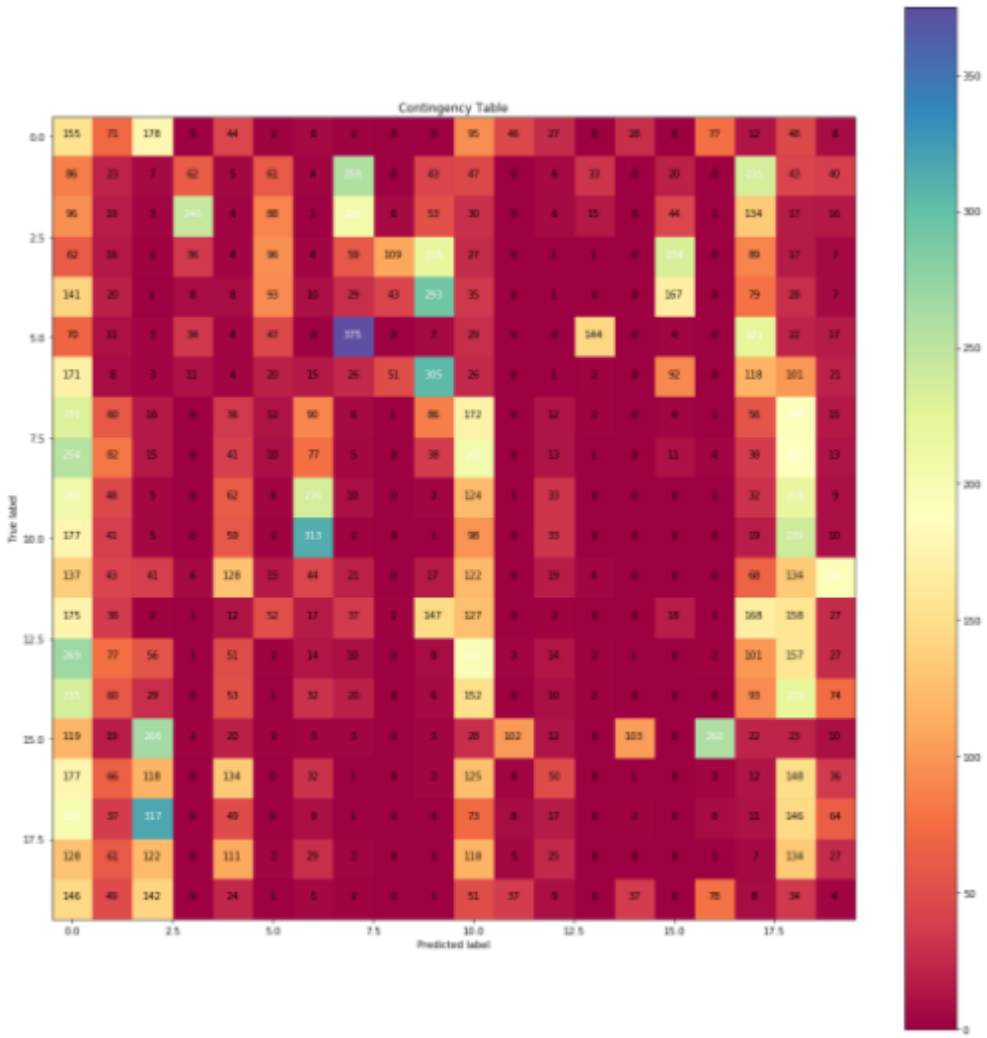

Model with 10 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	12.95s	148	0.230	0.259	0.244	0.072	0.227

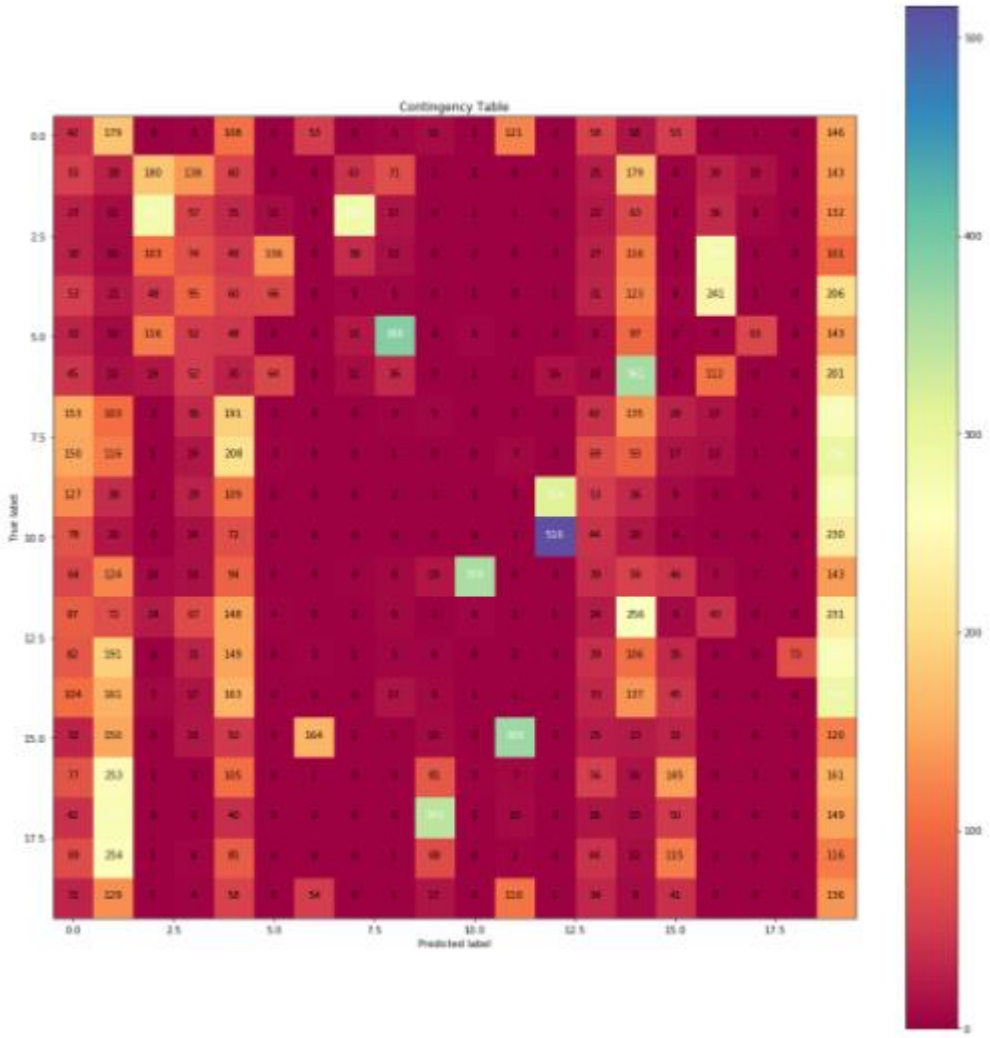

Model with 20 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	11.56s	334	0.264	0.299	0.280	0.081	0.261

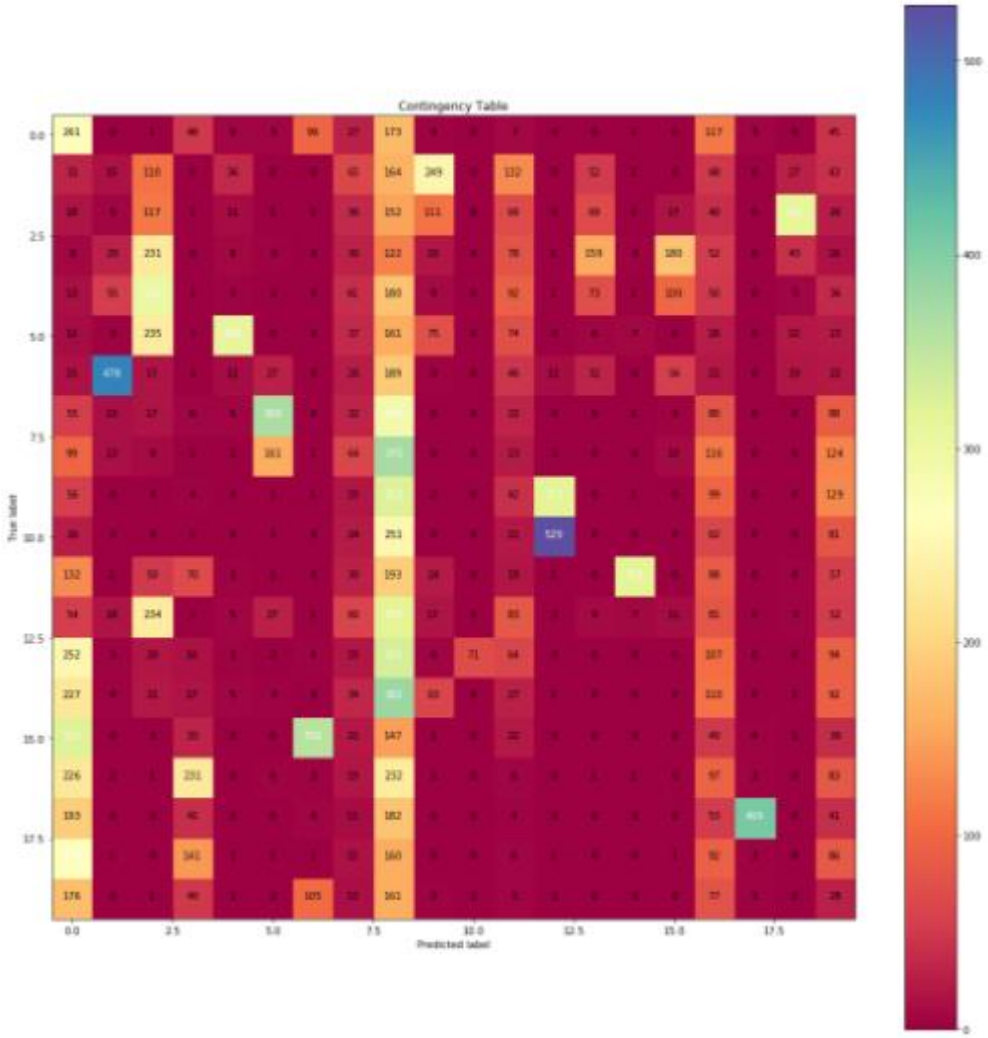

Model with 50 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	15.05s	860	0.253	0.315	0.281	0.057	0.250

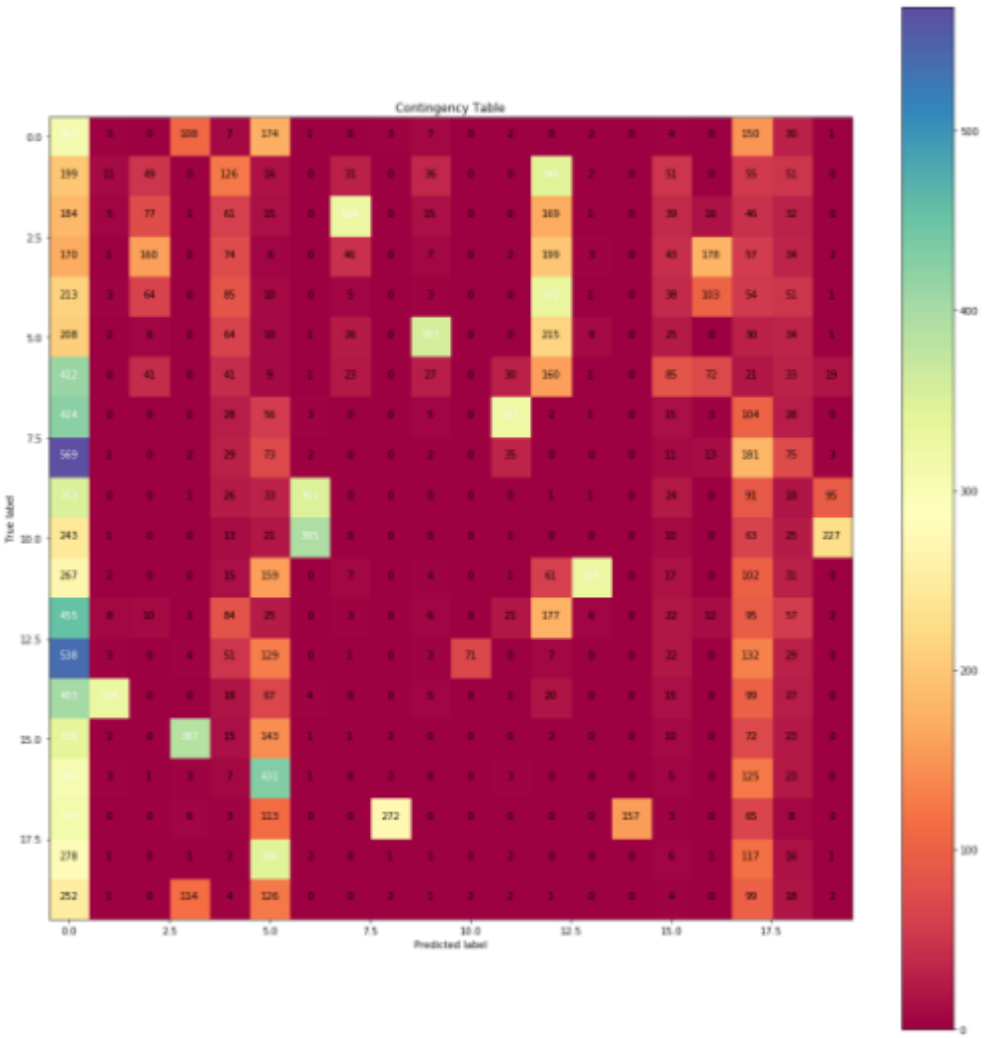

Model with 100 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	21.82s	1560	0.248	0.326	0.282	0.052	0.246

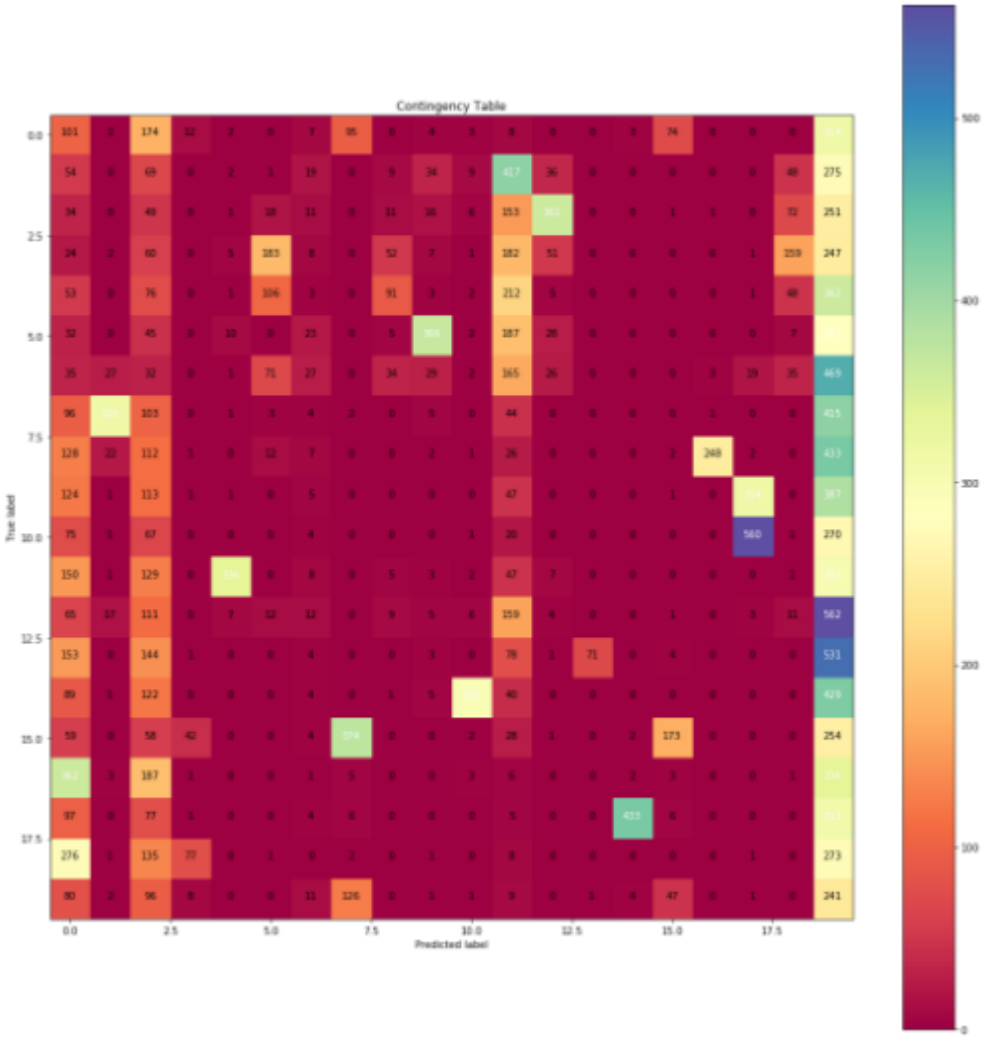

Model with 300 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	42.33s	3461	0.261	0.354	0.300	0.053	0.259

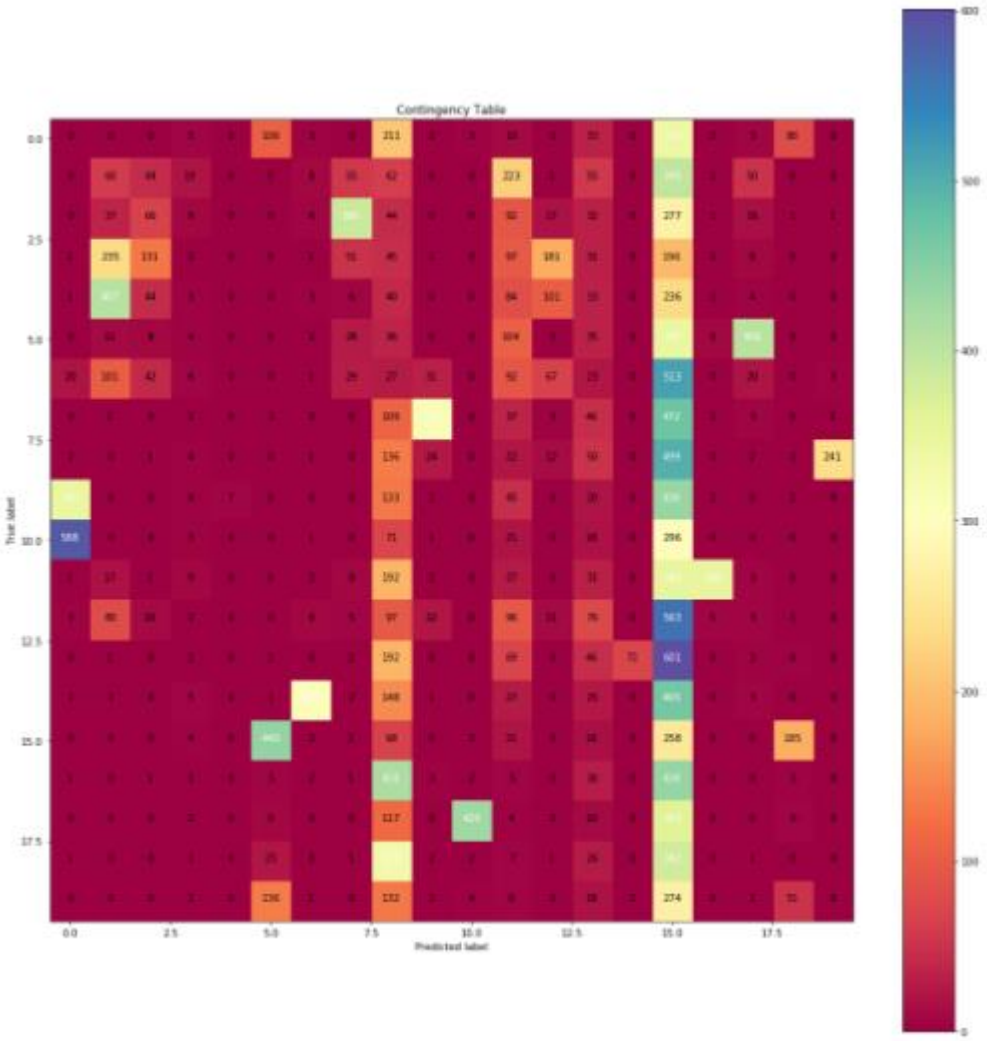
Contingency Table
True label
Predicted label
0.0
2.5
5.0
7.5
10.0
12.5
15.0
17.5
500
400
300
200
100
0

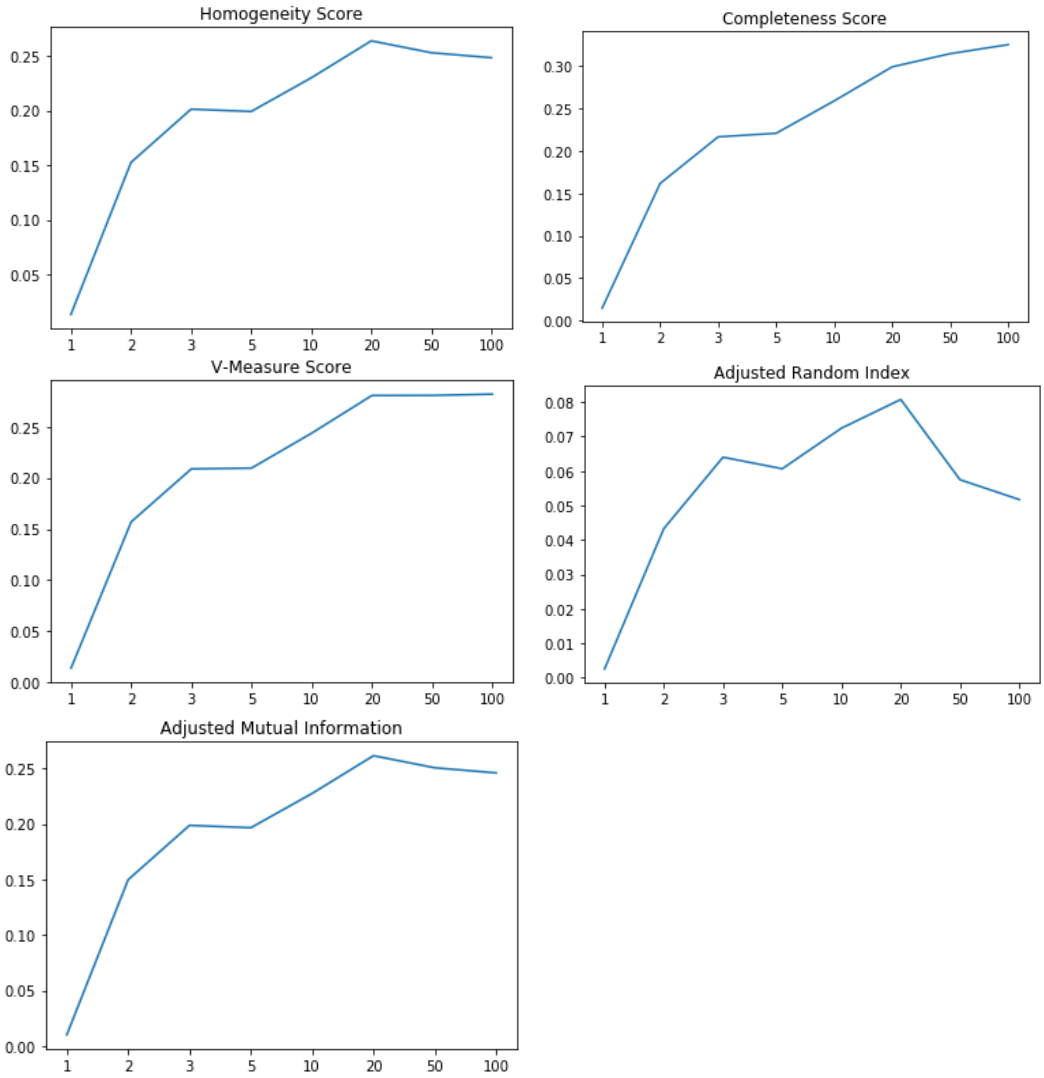

As we can observe from the contingency matrix as well as the 5 measures above, the best r value for LSI is 20. i.e, 20 components are best suitable to represent each document and cluster with maximum accuracy. Although, the graph for completeness score and V-Measure are somewhat monotonic in nature, we consider all the measures and decide 20 to be the best value.

b. NMF (Non-Negative Matrix Factorization):
After performing NMF on the TF-IDF matrix, following are the obtained results:

Model with 1 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	4.94s	0	0.013	0.015	0.014	0.003	0.010

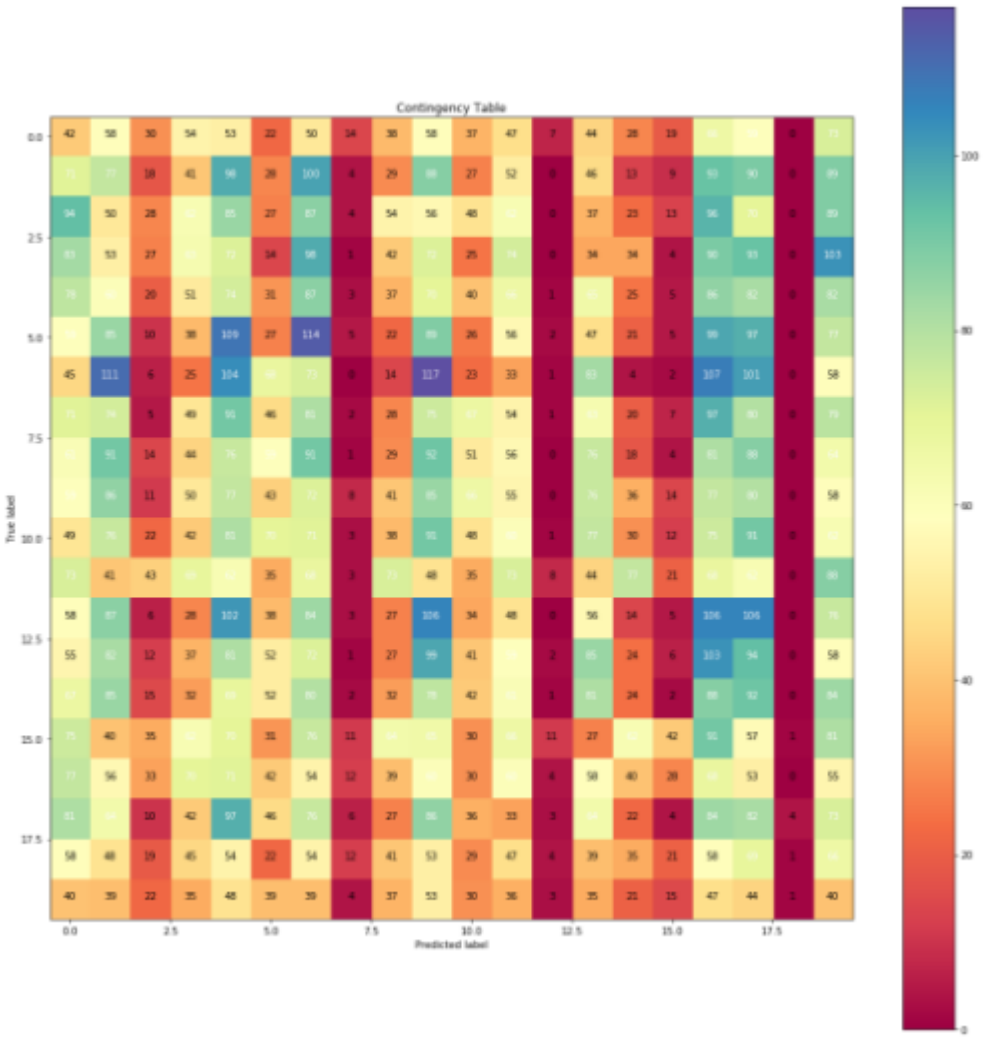

Model with 2 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	10.66s	0	0.143	0.154	0.149	0.041	0.141

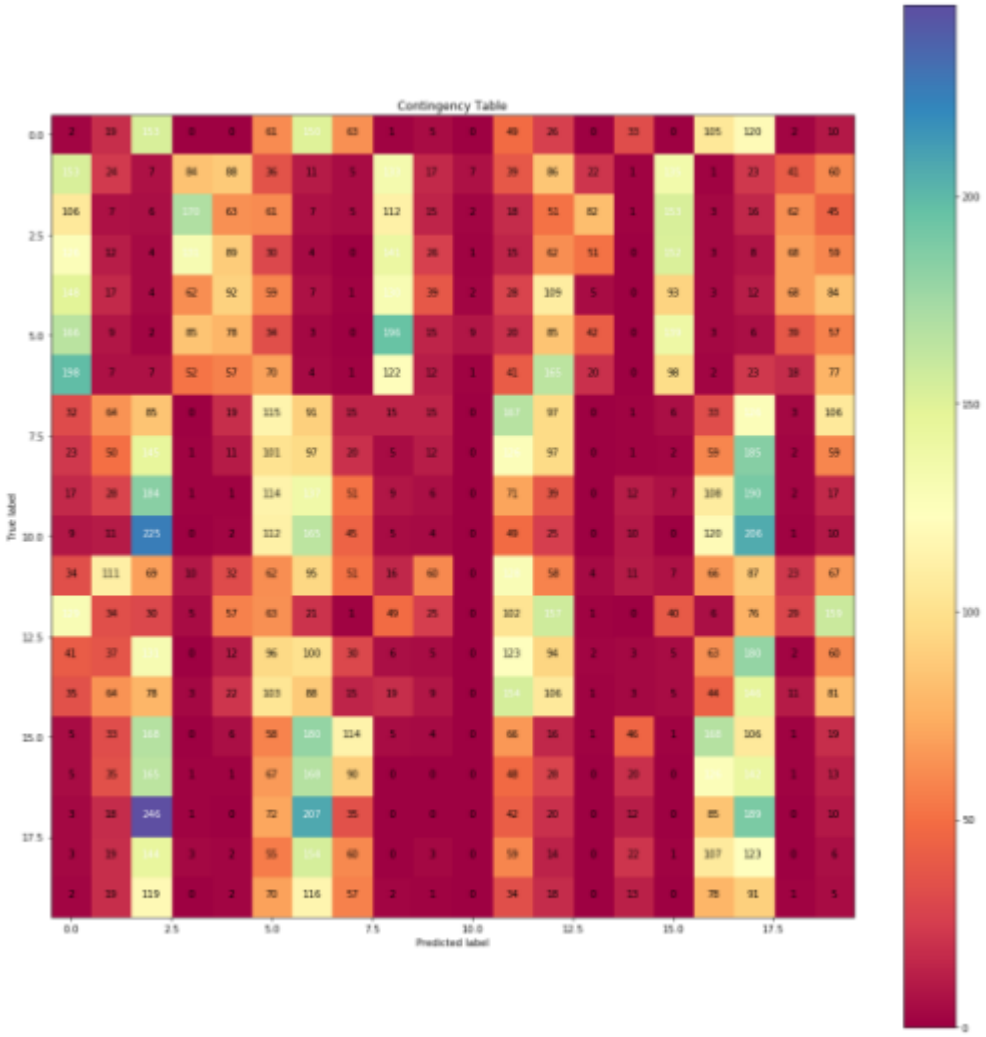

Model with 3 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	8.95s	0	0.181	0.199	0.190	0.052	0.178

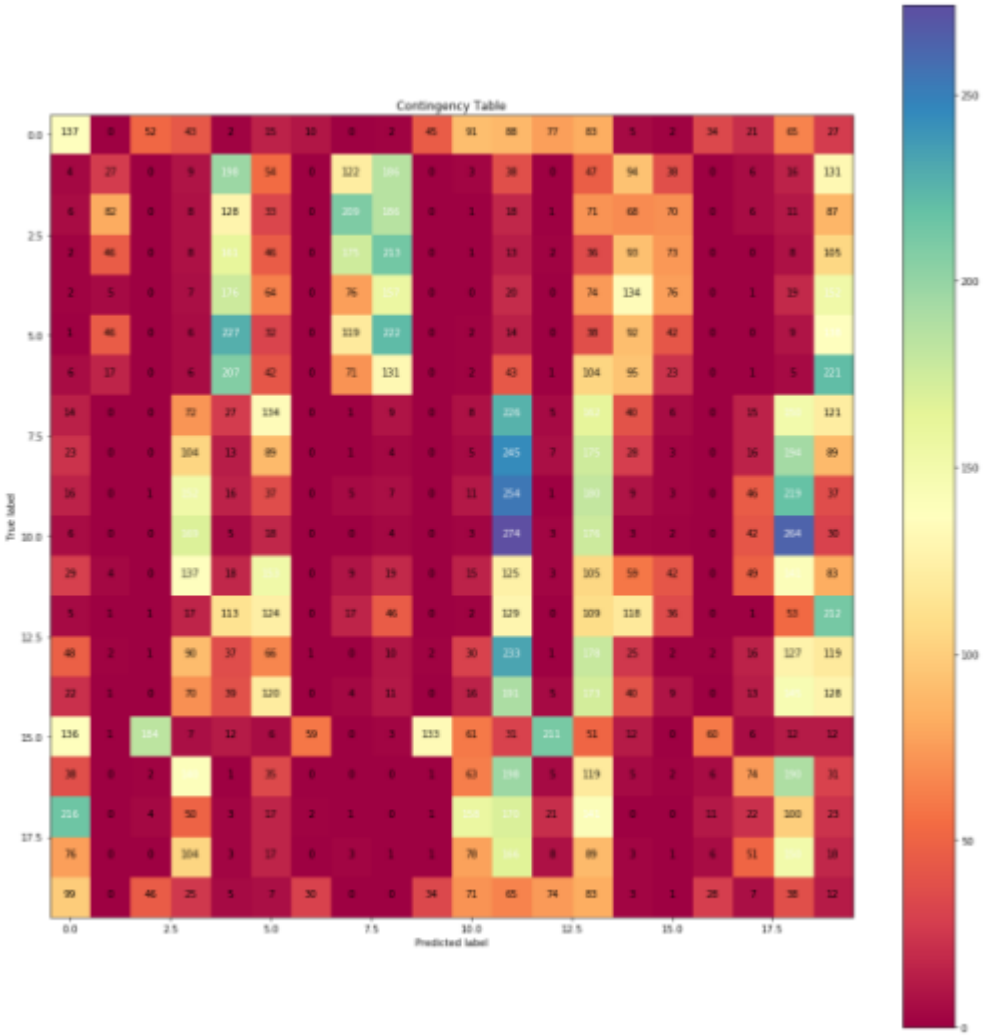

Model with 5 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	9.11s	1	0.174	0.191	0.182	0.049	0.171

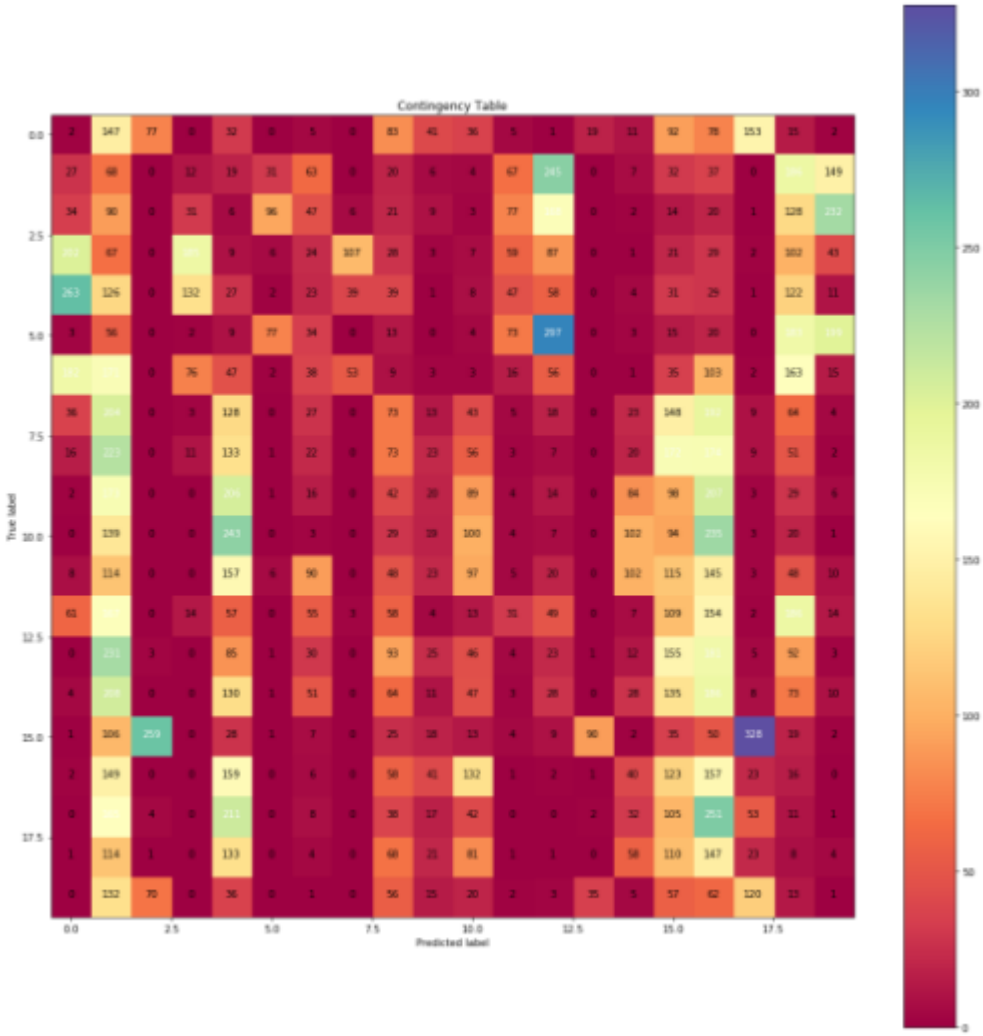

Model with 10 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	6.06s	6	0.214	0.248	0.230	0.057	0.212

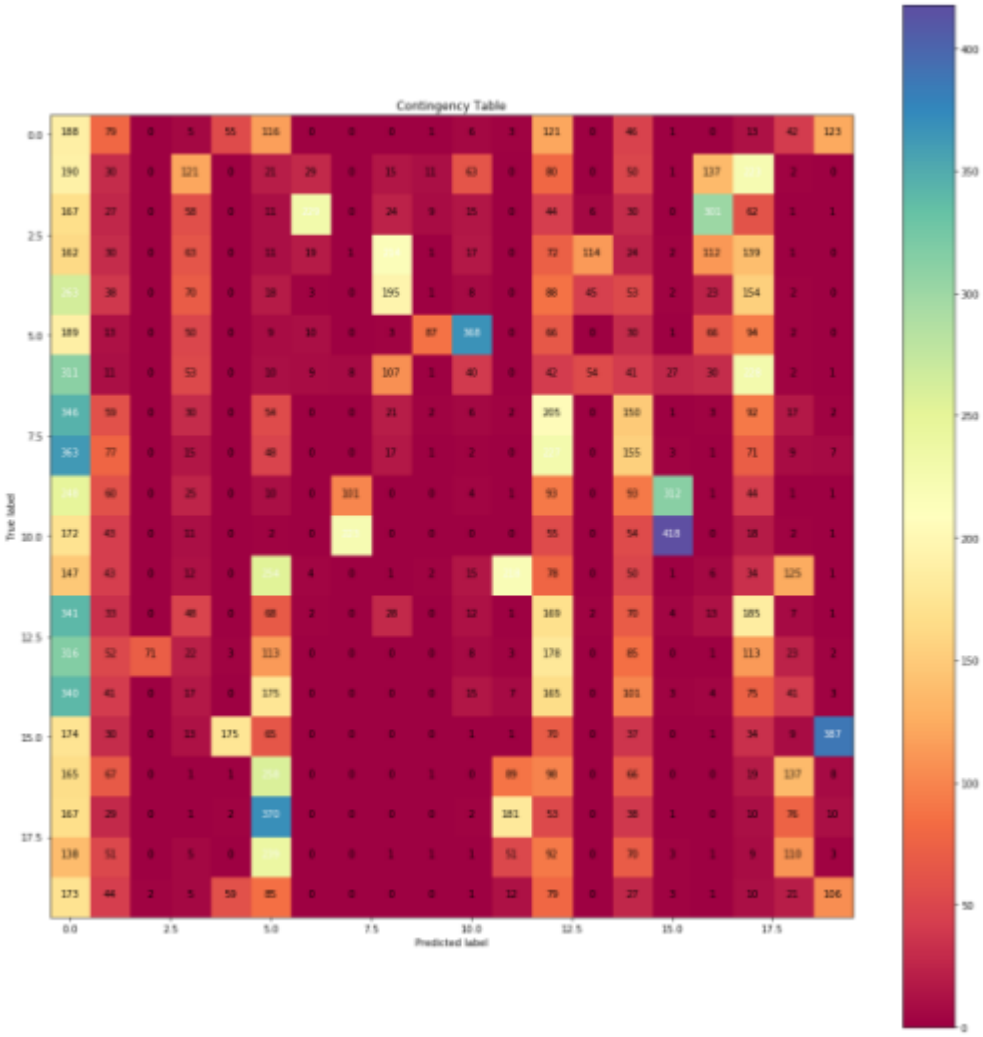

Model with 20 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	7.28s	19	0.227	0.281	0.251	0.049	0.225

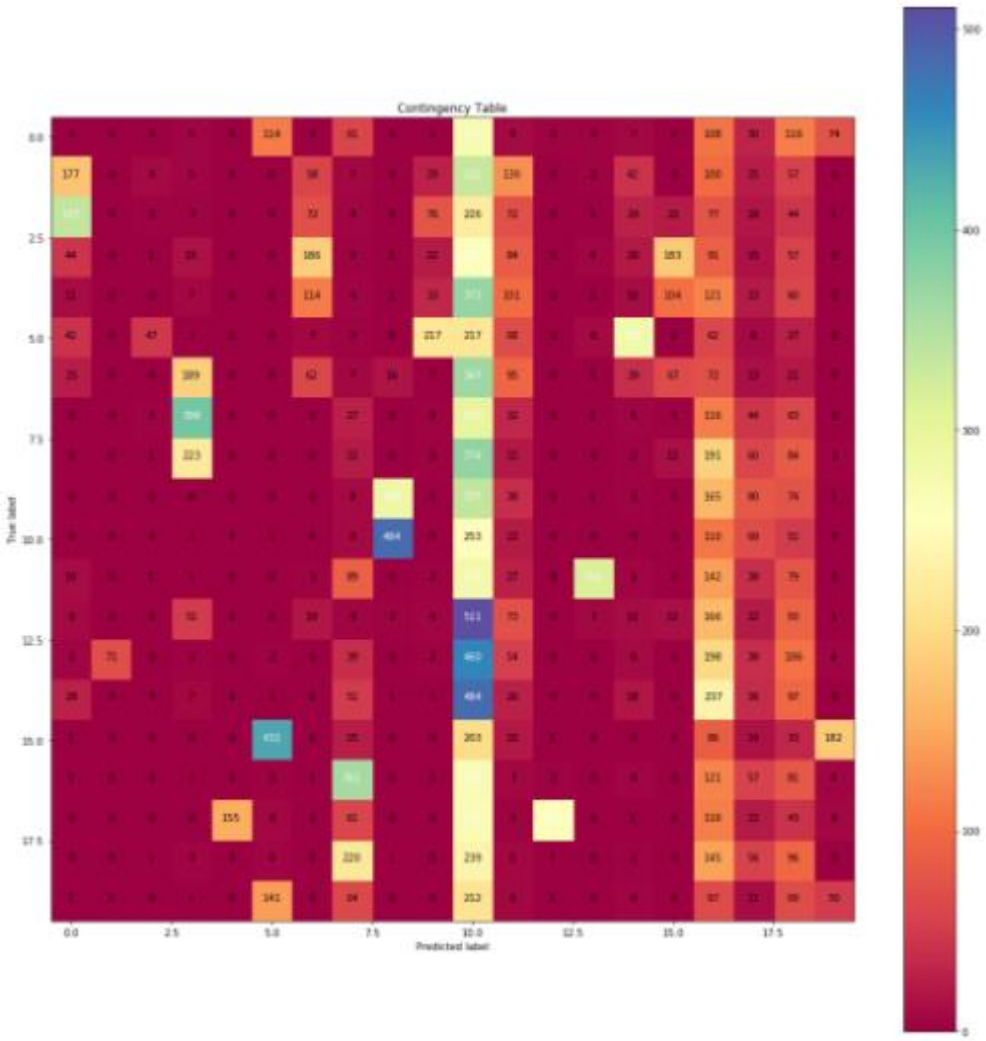

Model with 50 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	11.42s	70	0.179	0.258	0.211	0.033	0.176

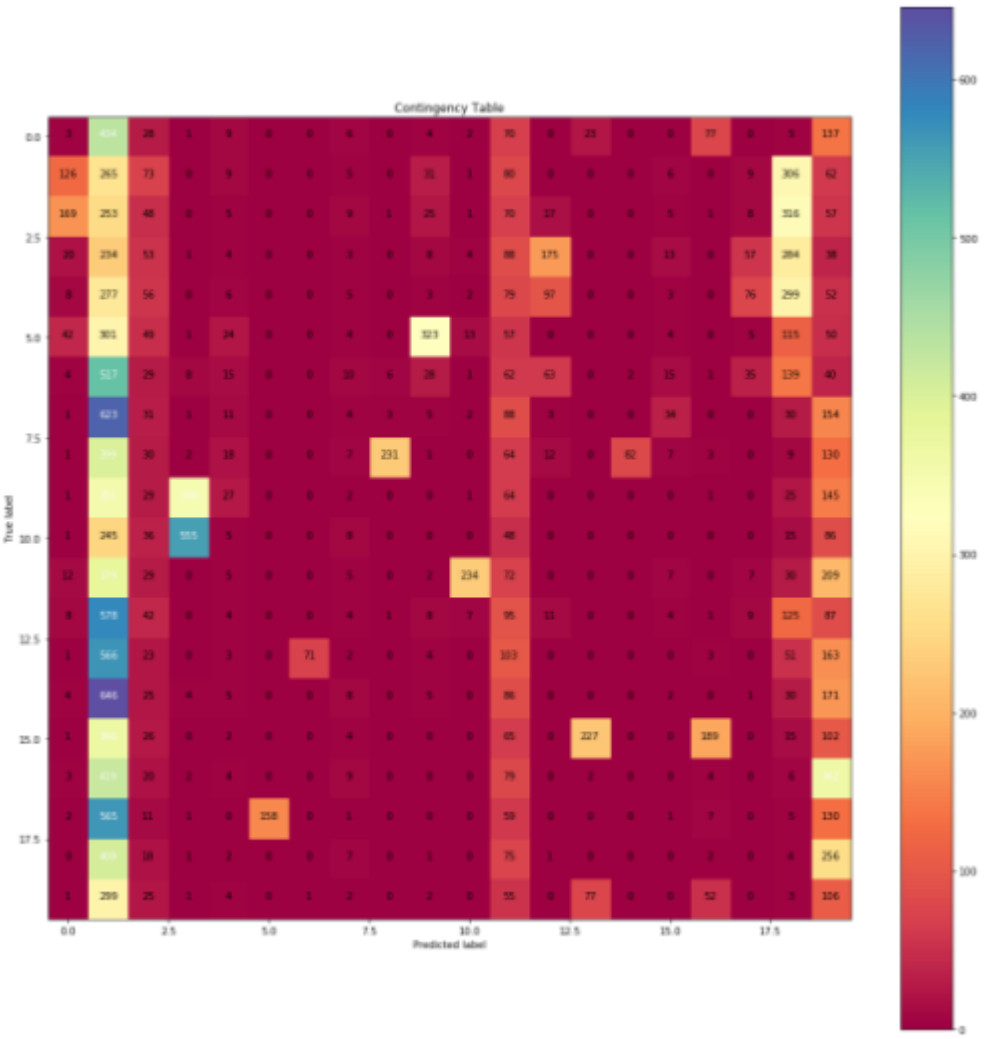

Model with 100 components:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	15.94s	158	0.102	0.149	0.121	0.016	0.099

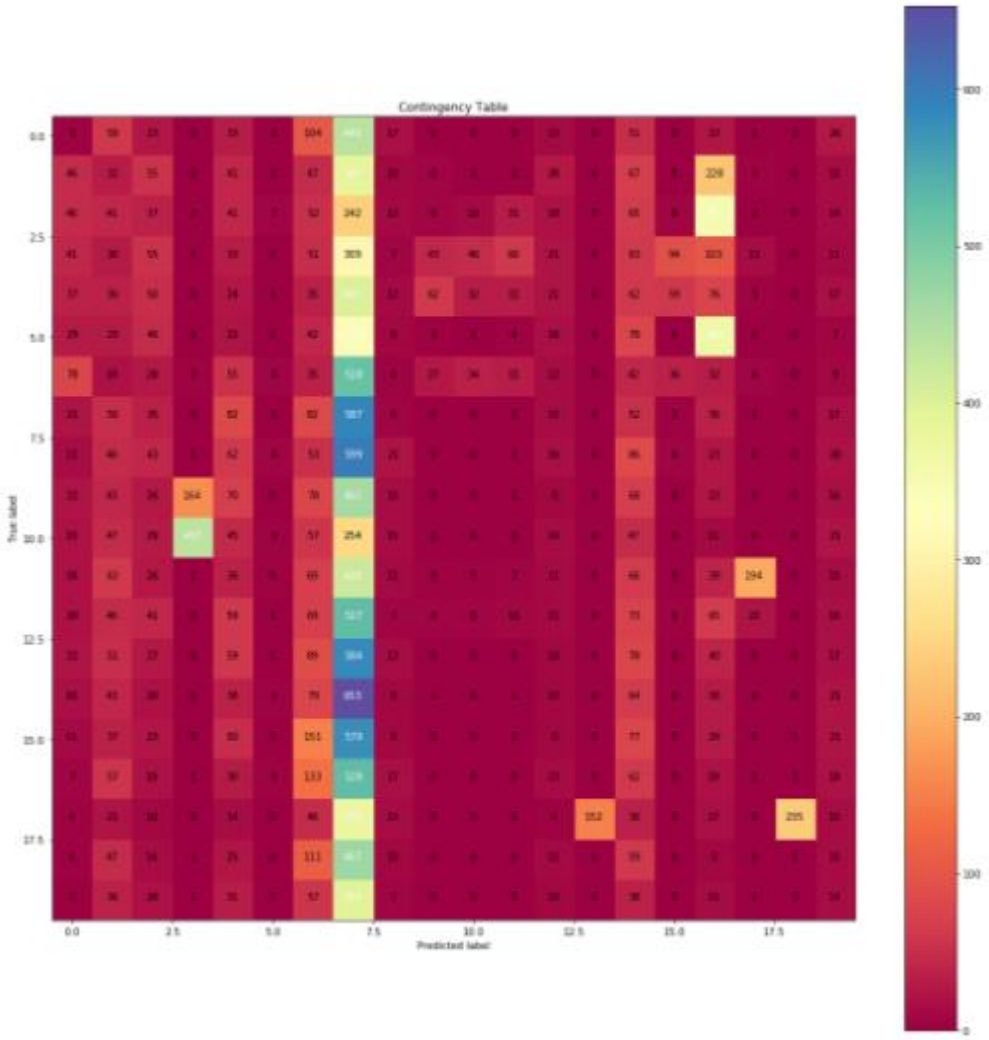
Contingency Table
True label
Predicted label
0.0
2.5
5.0
7.5
10.0
12.5
15.0
17.5
600
500
400
300
200
100
0

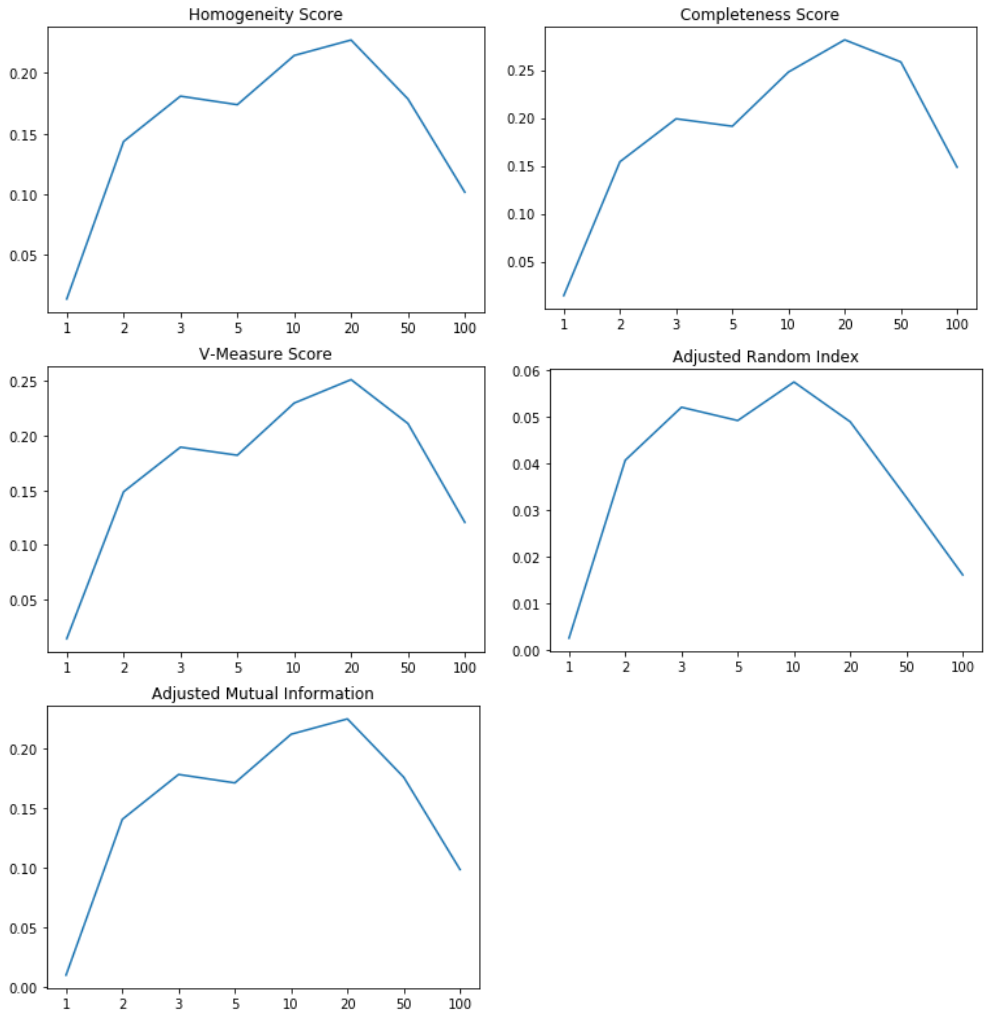

As we can observe, the best r value is at r = 20. i.e, only 20 components/features are more than enough to best cluster the documents in two classes with maximum accuracy.

4. Visualize the performance:
    a. We found that,

    r=5 works best in case of LSI, and
    r=2 works best in case of NMF.
    Hence, following are the clusters with their decision boundary for the two methods:

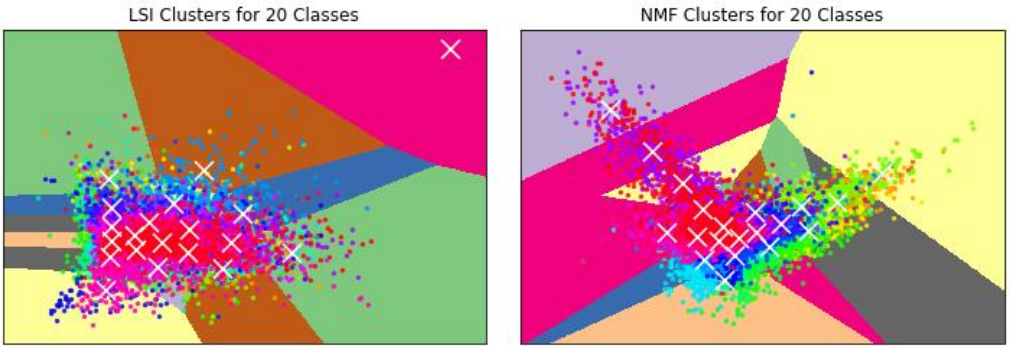

    b. 3 Methods:

i. Normalizing:
We performed scaling using sklearn “scale” to get unit variance and then normalized the data. Following are the results for LSI and NMF:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	11.56s	10	0.243	0.283	0.262	0.067	0.241

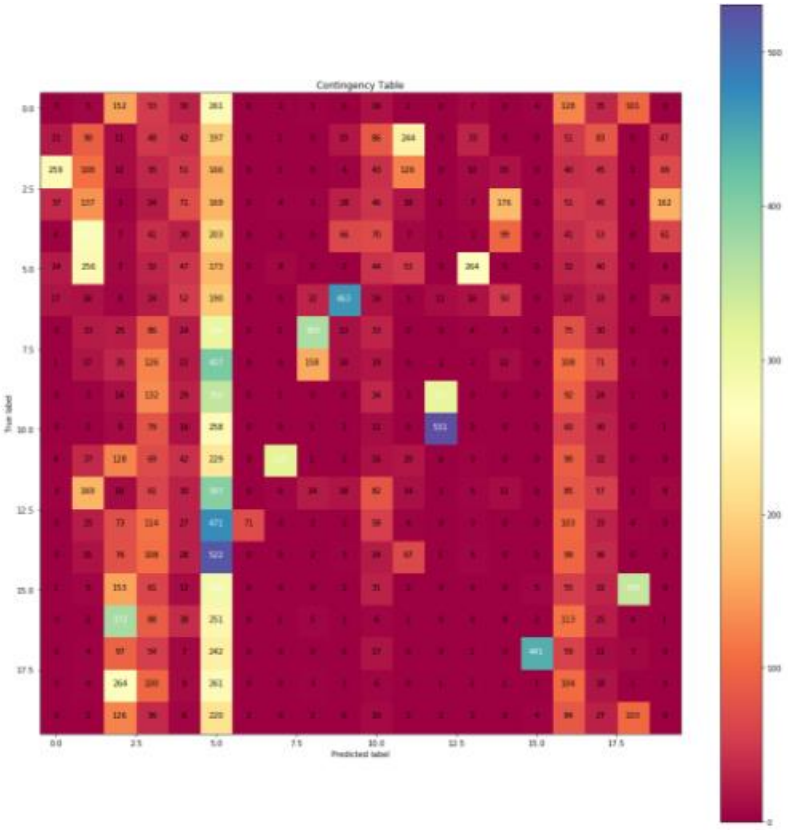

LSI Clusters afer Normalizing

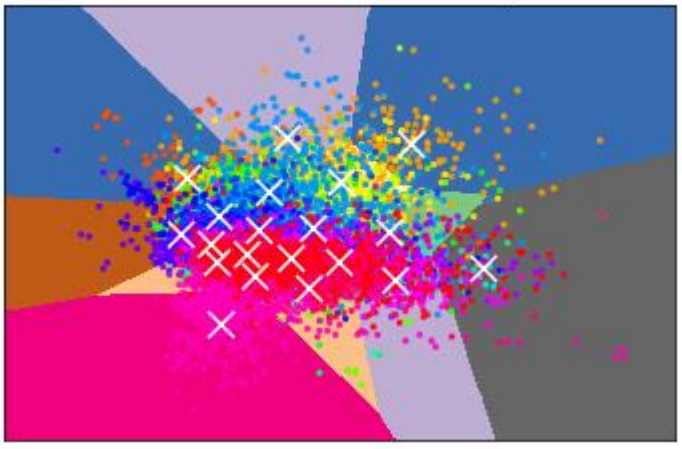

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	10.31s	7	0.243	0.303	0.270	0.053	0.241

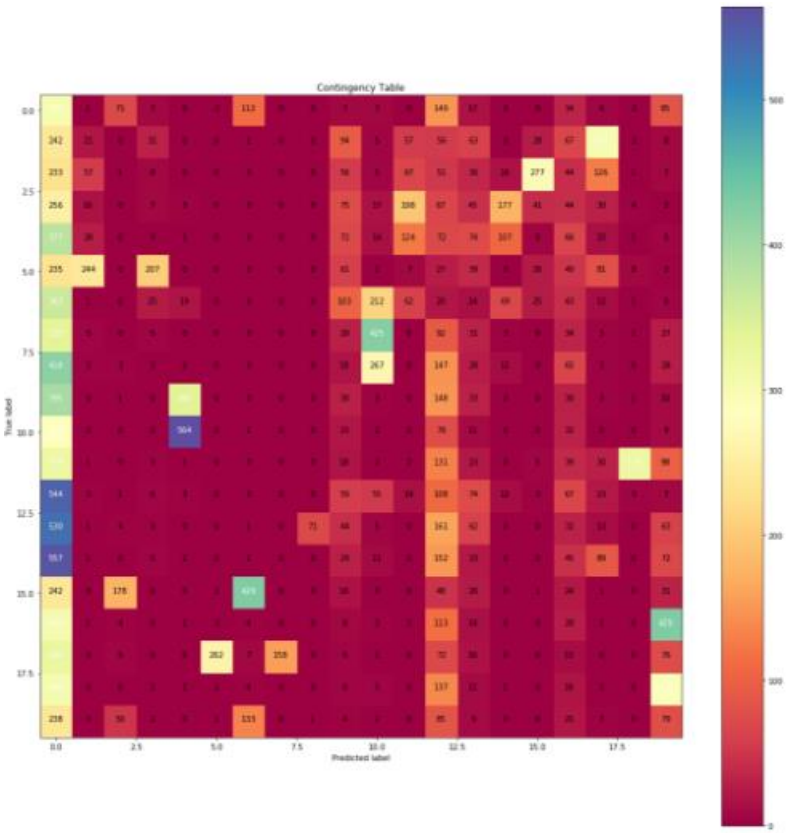

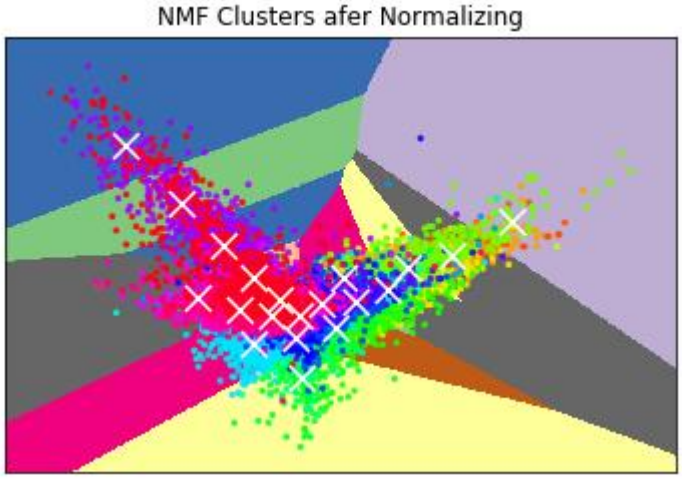

ii. Logarithmic Transformation:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	18.65s	253625	0.302	0.304	0.303	0.156	0.299

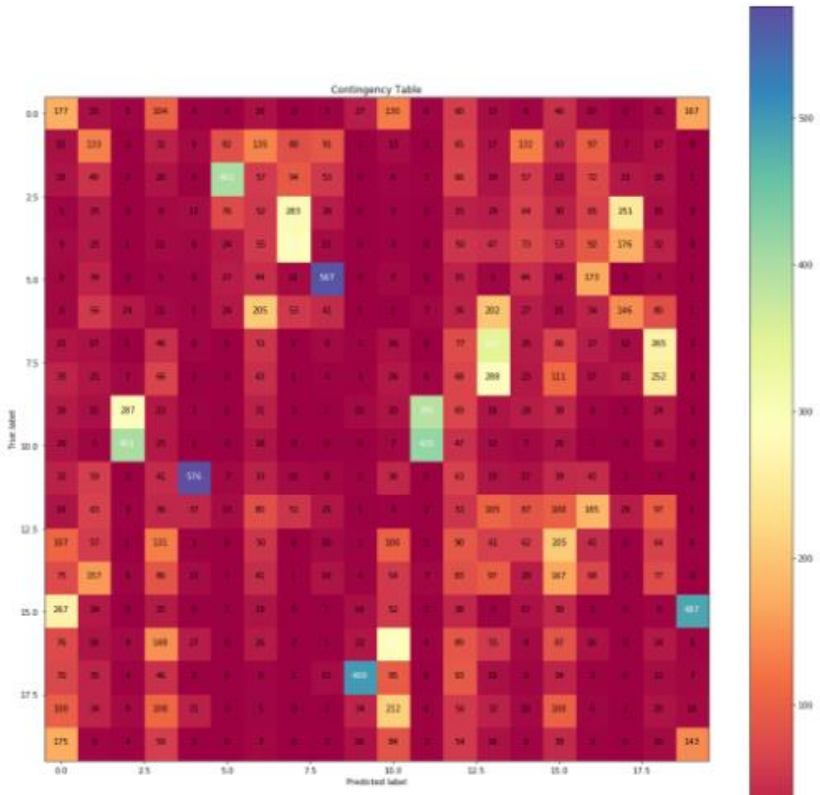

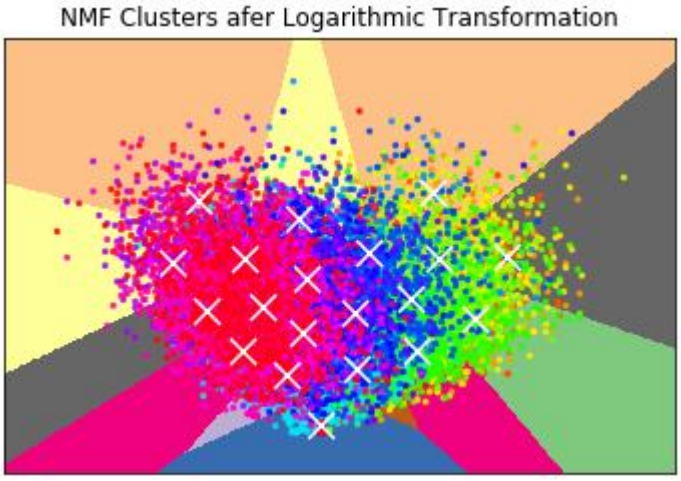

iii. Combination

1. Normalization and then Log Transformation
   We perform normalization and then apply log transformation on NMF data to get the following results:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	12.30s	11	0.325	0.330	0.327	0.186	0.323

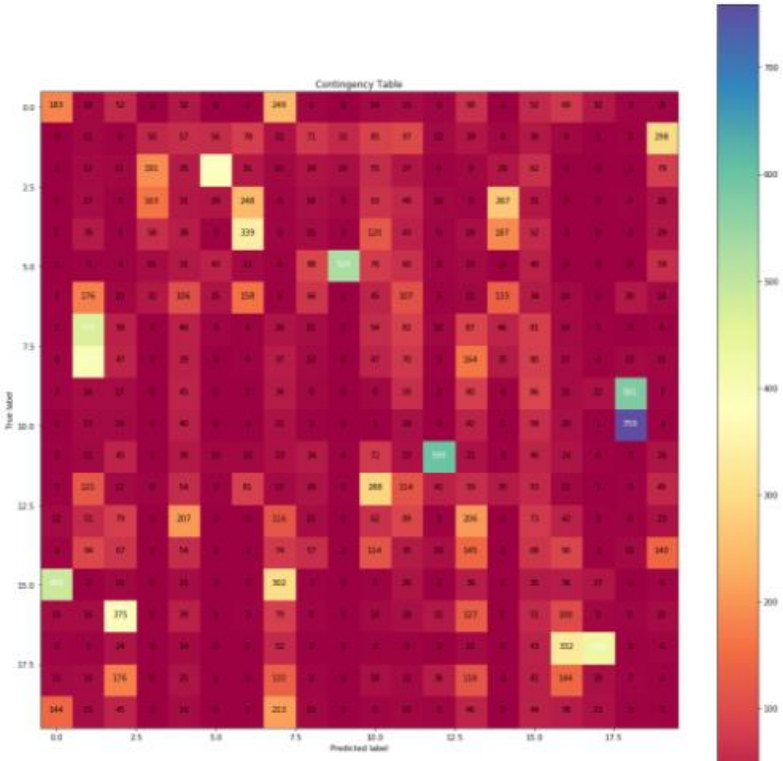

NMF Clusters afer Logarithmic Transformation and Normalizing

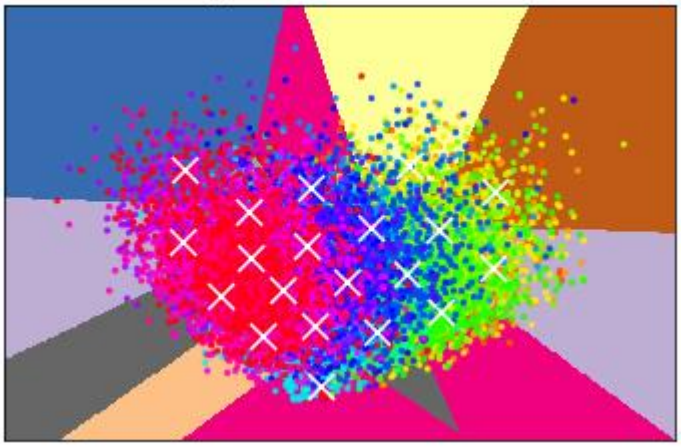

2. Log Transformation and then Normalization
   We apply log transformation on NMF data and then perform normalization to get the following results:

init	time	inertia	homo	comp	v-meas	ARI	AMI
k-means++	16.37s	176083	0.313	0.315	0.314	0.161	0.311

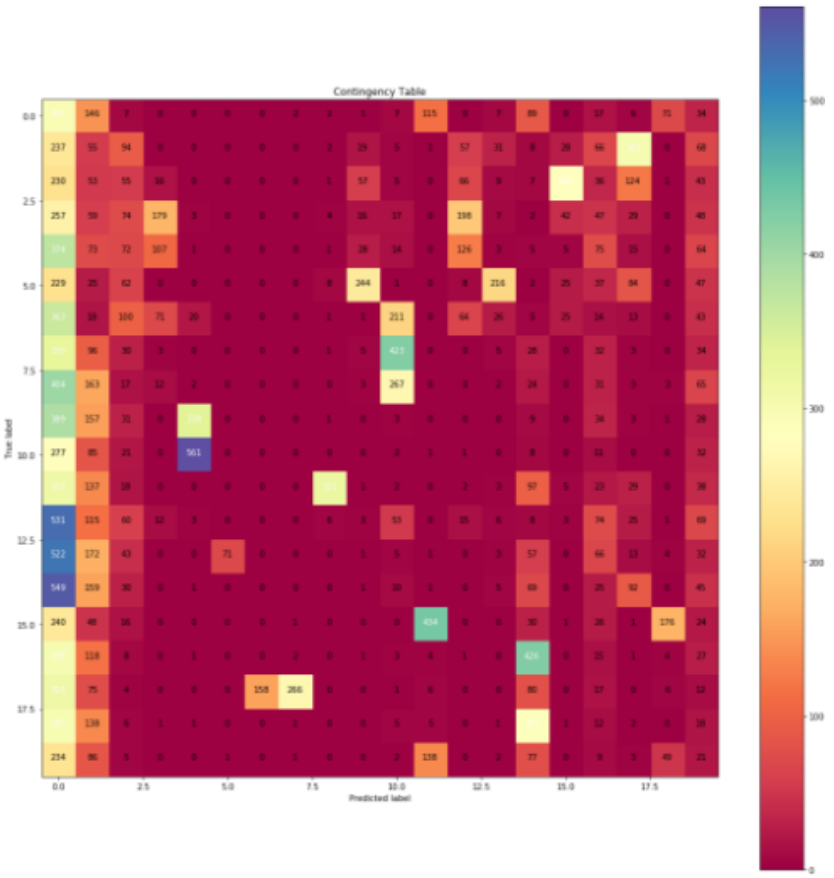

NMF Clusters afer Normalizing and Logarithmic Transformation